



hochschule mannheim

**Conception, Implementation, and Evaluation  
of a Proof of Concept of a Modular Proxy  
Application for Testing Internet of Things  
Applications**

Moritz Laurin Thomas

Master Thesis

for the acquisition of the academic degree Master of Science (M.Sc.)

Course of Studies: Computer Science

Department of Computer Science

University of Applied Sciences Mannheim

31.05.2021

Tutors

Prof. Dr. Thomas Specht, Hochschule Mannheim

Pierre-Alain Mouy, M.Sc., NVISO GmbH

**Thomas, Moritz Laurin:**

Conception, Implementation, and Evaluation of a Proof of Concept of a Modular Proxy Application for Testing Internet of Things Applications / Moritz Laurin Thomas. –

Master Thesis, Mannheim: University of Applied Sciences Mannheim, 2020. 88 pages.

**Thomas, Moritz Laurin:**

Konzeption, Implementierung und Evaluation eines Machbarkeitsnachweises eines modularen Proxys zum Testen von Anwendungen im Internet der Dinge / Moritz Laurin Thomas. –

Master-Thesis, Mannheim: Hochschule Mannheim, 2020. 88 Seiten.

## **Erklärung**

Hiermit erkläre ich, dass ich die vorliegende Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe.

Ich bin damit einverstanden, dass meine Arbeit veröffentlicht wird, d. h. dass die Arbeit elektronisch gespeichert, in andere Formate konvertiert, auf den Servern der Hochschule Mannheim öffentlich zugänglich gemacht und über das Internet verbreitet werden darf.

Mannheim, 31.05.2021

Moritz Laurin Thomas



# Abstract

## ***Conception, Implementation, and Evaluation of a Proof of Concept of a Modular Proxy Application for Testing Internet of Things Applications***

Today, more and more formerly analogue physical entities are now being digitized and connected to the internet, adding to the “Internet of Things”. However, the wide variety in appliances poses a potentially wide attack surface for malicious actors. To address this risk that these so-called *smart* devices pose to parties that employ them, security researchers and penetration testers examine and test their security implementation. The need arises for a modular proxy application that allows to test the heterogeneous landscape of communication protocols used in IoT applications. The goal of this thesis is to conceptualize a design for such an application, realize a prototypic implementation thereof and evaluate its usefulness. Quantitative results are a documentation of the problem space, an abstract design concept and sets of development challenges and lessons learned.

## ***Konzeption, Implementierung und Evaluation eines Machbarkeitsnachweises eines modularen Proxys zum Testen von Anwendungen im Internet der Dinge***

Als Konsequenz der voranschreitenden Digitalisierung werden ehemals analoge Geräte zunehmend digitalisiert und somit Teil des “Internets der Dinge” (“Internet of Things”, IoT). Dabei stellt jedoch die große Bandbreite an Anwendungen eine potenziell große Angriffsfläche für Angreifer dar. Um diesem Risiko, das die sogenannten “smarten” Anwendungen gegenüber ihren Betreibern darstellen, zu begegnen, untersuchen und überprüfen Sicherheitsforscher und Penetrationstester deren Sicherheitsarchitekturen. Daraus erwächst ein Bedarf an einer modularen Proxy-Anwendung, die sie dabei unterstützt, die heterogene Verwendung von Kommunikationsprotokollen in IoT-Anwendungen zu beherrschen. Ziel dieser Arbeit ist die Konzeption eines Softwareentwurfs für eine solche Anwendung und deren prototypische Umsetzung sowie eine Bewertung ihrer Nützlichkeit. Quantitative Ergebnisse sind die Dokumentation der Problemstellung, ein abstraktes Entwurfskonzept und eine Reihe von Herausforderungen bei der Entwicklung sowie daraus gewonnene Erkenntnisse.



# Acknowledgements

In this moment of accomplishment in these troublesome times, I would like to thank my supervisor Prof. Dr. Thomas Specht for convincing me of the value of my work and its results and thereby encouraging me to finish this thesis at a point of resignation. Furthermore, his extensive and constructive feedback was of immense value.

I took great inspiration from the conversations I had with my supervisor Pierre-Alain Mouy. His experience, motivating character and solution-oriented mindset positively influenced the design concept phases in particular. I would like to thank him for introducing me to the idea of working on this thesis and inviting me to work on it at - and as part of - *NVISO*.

This thesis would lack important input from professionals if it was not for the openness of the volunteers at *NVISO* which took their time and participated in interviews. I would like to express my gratitude to Patrick Eisenschmidt, Oliver Nettinger, Cédric Bassem, Théo Rigas, Pierre-Alain Mouy, and Jonah Bellemans for the time, effort and interest they put into these interviews.

Further, I want to thank Patrick Eisenschmidt and Adrian Gast for proof-reading this thesis and providing valuable feedback.

Finally, I cannot begin to express my thanks to my wife Marie-Darline Thomas for her sincere, ongoing and unfailing love and support. I deeply appreciate her belief in me, her encouragement and her being there when I need her.





# Contents

|  |           |
|--|-----------|
| <b>1. Introduction</b>                               | <b>1</b>  |
| 1.1. Motivation . . . . .                            | 1         |
| 1.2. Purpose and Structure of the Thesis . . . . .   | 2         |
| <b>2. Related Work</b>                               | <b>3</b>  |
| <b>3. Theoretical Background</b>                     | <b>5</b>  |
| 3.1. Design Patterns . . . . .                       | 5         |
| 3.1.1. Pipeline/Pipes and Filters Pattern . . . . .  | 5         |
| 3.1.2. Abstract Factory/Kit Pattern . . . . .        | 6         |
| 3.1.3. Publish-Subscribe/Observer Pattern . . . . .  | 7         |
| 3.2. Computer Communication: The OSI-Model . . . . . | 7         |
| 3.3. Internet of Things . . . . .                    | 9         |
| 3.3.1. Fields of Use . . . . .                       | 9         |
| 3.3.2. Common Protocols . . . . .                    | 10        |
| 3.4. Information Security . . . . .                  | 12        |
| 3.4.1. The CIA Triad . . . . .                       | 13        |
| 3.4.2. Methodology . . . . .                         | 13        |
| 3.4.3. Tools . . . . .                               | 15        |
| <b>4. Procedure</b>                                  | <b>19</b> |
| <b>5. Understanding the Problem Space</b>            | <b>21</b> |
| 5.1. Prototypical Implementation . . . . .           | 21        |
| 5.1.1. Example Scenarios . . . . .                   | 22        |
| 5.1.2. Requirements . . . . .                        | 27        |
| 5.1.3. Design . . . . .                              | 29        |
| 5.1.4. Testing . . . . .                             | 32        |
| 5.1.5. Implementation . . . . .                      | 35        |
| 5.1.6. Insights Gained . . . . .                     | 38        |
| 5.2. Interviewing Experts for Insights . . . . .     | 39        |
| 5.2.1. Interview Guideline . . . . .                 | 40        |
| 5.2.2. Conducting Interviews . . . . .               | 40        |
| 5.2.3. Interview Analysis . . . . .                  | 41        |

|  |             |
|--|-------------|
| 5.3. Analysis of Existing Software . . . . .           | 45          |
| <b>6. Conceptual Design</b>                            | <b>47</b>   |
| 6.1. Design #1: Monolithic Proxy Application . . . . . | 47          |
| 6.1.1. High-level Overview . . . . .                   | 47          |
| 6.1.2. State-Machines . . . . .                        | 51          |
| 6.1.3. Gateway . . . . .                               | 52          |
| 6.1.4. Pipes . . . . .                                 | 53          |
| 6.1.5. Messages . . . . .                              | 55          |
| 6.2. Design #2: Distributed Proxy Services . . . . .   | 56          |
| 6.2.1. Overview . . . . .                              | 57          |
| 6.2.2. Registry and Units . . . . .                    | 57          |
| 6.2.3. PipelineRepository . . . . .                    | 58          |
| 6.2.4. State of the Design Concept . . . . .           | 59          |
| 6.2.5. Comparison of Both Designs . . . . .            | 59          |
| <b>7. Implementing the Modular Proxy Application</b>   | <b>61</b>   |
| 7.1. Goals and Constraints . . . . .                   | 61          |
| 7.2. Tool Selection . . . . .                          | 62          |
| 7.3. Individual Components . . . . .                   | 63          |
| 7.3.1. Gateways . . . . .                              | 63          |
| 7.3.2. Encoders . . . . .                              | 64          |
| 7.3.3. Scripting . . . . .                             | 65          |
| 7.3.4. Configuration Parsing and Building . . . . .    | 67          |
| <b>8. Evaluation and Lessons Learned</b>               | <b>71</b>   |
| 8.1. Quantitative Overview: Time Management . . . . .  | 71          |
| 8.1.1. Project Timeline . . . . .                      | 71          |
| 8.1.2. Development Challenges . . . . .                | 75          |
| 8.2. Qualitative Overview: Deliverables . . . . .      | 78          |
| 8.3. Lessons Learned . . . . .                         | 80          |
| <b>9. Summary</b>                                      | <b>83</b>   |
| 9.1. Design Concepts . . . . .                         | 83          |
| 9.2. Implementation . . . . .                          | 84          |
| <b>10. Conclusion</b>                                  | <b>87</b>   |
| 10.1. Outlook . . . . .                                | 88          |
| <b>List of Abbreviations</b>                           | <b>xi</b>   |
| <b>List of Figures</b>                                 | <b>xiii</b> |
| <b>List of Tables</b>                                  | <b>xvii</b> |
| <b>Listings</b>  | <b>xix</b>  |

|                               |              |
|-------------------------------|--------------|
| <b>Bibliography</b>           | <b>xxi</b>   |
| <b>Index</b>                  | <b>xxv</b>   |
| <b>A. Diagrams</b>            | <b>xxv</b>   |
| <b>B. Interview Guideline</b> | <b>xxvii</b> |



# Chapter 1

## Introduction

This chapter will introduce the underlying motivation of this thesis. Then, it will give an overview of this thesis' purpose and structure.

### 1.1. Motivation

Today scientific and industrial parties work on connecting physical entities such as machines, buildings and even humans to the internet by equipping them with digital sensors and actuators, referred to as “Internet of Things (IoT)”. While this progression promises many positive effects, such as simplifying tasks in our personal day-to-day life (“Smart Home” applications), monitoring our personal health (“eHealth”) and increasing efficiency and safety of industrial plants (“Industrial Internet of Things (IIoT)”, also referred to as “Industry 4.0”), it also yields the risk of introducing new attack-vectors to parts of our environment: “smart” devices used at home or at other sensitive places may introduce weak security implementations or faulty security design, resulting in private and personal data being available to parties interested in violating the privacy of one’s home (e.g. vacuum robots leaking information about the interior design of homes[46]) or conducting industrial espionage which is an acute threat [4, p. 14].

The diversity of both deployed smart devices and the internet services those devices are connected to, lead to the need and use of ever-increasing complex technologies used for communication, data storage and access management, further adding to potential attack-vectors of connected devices and distributed applications [19, p. 119]. This complexity and the sheer number of connected devices is actively being ex-

exploited by attackers today and the number of attacks on IoT devices is increasing [10].

There are security guidelines, best practices and innovative approaches for developing secure smart applications [19, p. 120][20, p.326-328], however testing such applications proves to be cumbersome: intercepting, dissecting, inspecting and manipulating the communication in these applications requires working on various abstraction layers. In order to evaluate the security of such applications, penetration testers often spend a considerable amount of time dissecting applications and setting up test-environments.

The goal of this thesis is to conceptualize, implement and evaluate a modular proxy application that aids in examination of the security implementations of IoT applications.

### 1.2. Purpose and Structure of the Thesis

This thesis is separated into nine chapters: chapter 2 will give an overview of and discuss related and previous work. After that, relevant fundamentals about computer networks, IoT applications and information security will be covered in chapter 3.

The chapters 5 to 8 describe the research and development process of the IoT proxy application in chronological order: the problem space of the application is shown and dissected in chapter 5, yielding essential insights into potential challenges and technical requirements. Building upon these, the conceptual design of the IoT proxy application is proposed in chapter 6. This includes the process of collecting, documenting and analysis of software requirements, describing the application's work context and designing a software architecture that complies with the aforementioned requirements. Subsequently, chapter 7 involves a prototypical implementation of the above-said software concept, focusing on the goals and constraints of the implementation, the tools and frameworks used and the realization of core components of the application. The resulting implementation and the project itself are then analysed in an evaluation in chapter 8.

The thesis ends with a summary of all results produced and conclusions drawn from the work on this project.

## Chapter 2

### Related Work

As part of their master’s thesis, Bellemans conducted a study in 2020 that evaluated the security and privacy implementations of fifteen “*smart*” devices from a wide price range available on the market at the time. They performed automated analyses and requested data access from manufacturers [7]. The thesis showed that the devices made use of a variety of both standardized and proprietary transport and application protocols. It also found severe flaws in the devices’ compliance to General Data Protection Regulation (GDPR): about one third of the devices’ manufacturers did not reply to GDPR requests at all, however Bellemans noted that the COVID-19 pandemic may have had an impact on their data access requests. The thesis suggests that the introduction of a quality label that guarantees appropriate implementation of security and privacy aspects could prove beneficial for customers.

In 2017, Aphorpe et al. presented a three stage strategy to examine metadata of network traffic of four smart devices [3]. By monitoring the devices’ traffic, they showed that even though the communication between the devices and their corresponding internet servers were encrypted, passive observers could deduce information about users’ behaviour by identification of the destination server and analysis of the rate of traffic being sent. A noteworthy aspect of their work is that they performed this analysis from an Internet Service Provider (ISP)’s point of view, exclusively examining metadata of the communication that took place. The strategy described in the paper consists of the following (greatly simplified) steps:

1. Identifying communication streams of individual devices (e.g. by examining the TCP packets’ destination IPs).

## 2. Related Work

---

2. Associating the streams with specific device models (e.g. by performing reverse-look ups of the aforementioned IPs).
3. Analysing traffic rates (presuming that traffic is generated upon taking measures).

Apthorpe et al. conclude that their strategy works well on inferring behaviour from regular internet traffic of smart devices, however they suppose that shaping traffic or making use of proxies (that effectively mask the destination IPs) could be effective counter-measures. It is safe to assume that regular smart home setups do not make use of proxies or traffic shaping though, thus being vulnerable to this kind of attack.



## Chapter 3

# Theoretical Background

This chapter provides an overview of the technologies and concepts referred to in subsequent chapters. Starting with section 3.2, essential concepts of computer communication in networks will be presented and examined, detailing the concept of network layers. Building upon these fundamentals, section 3.3 introduces the fields of use of IoT applications popular protocols they make use of. After that, section 3.4 will provide insights into relevant concepts and the practices used and applied in information security. It information security properties, methodology and tools used in information security.

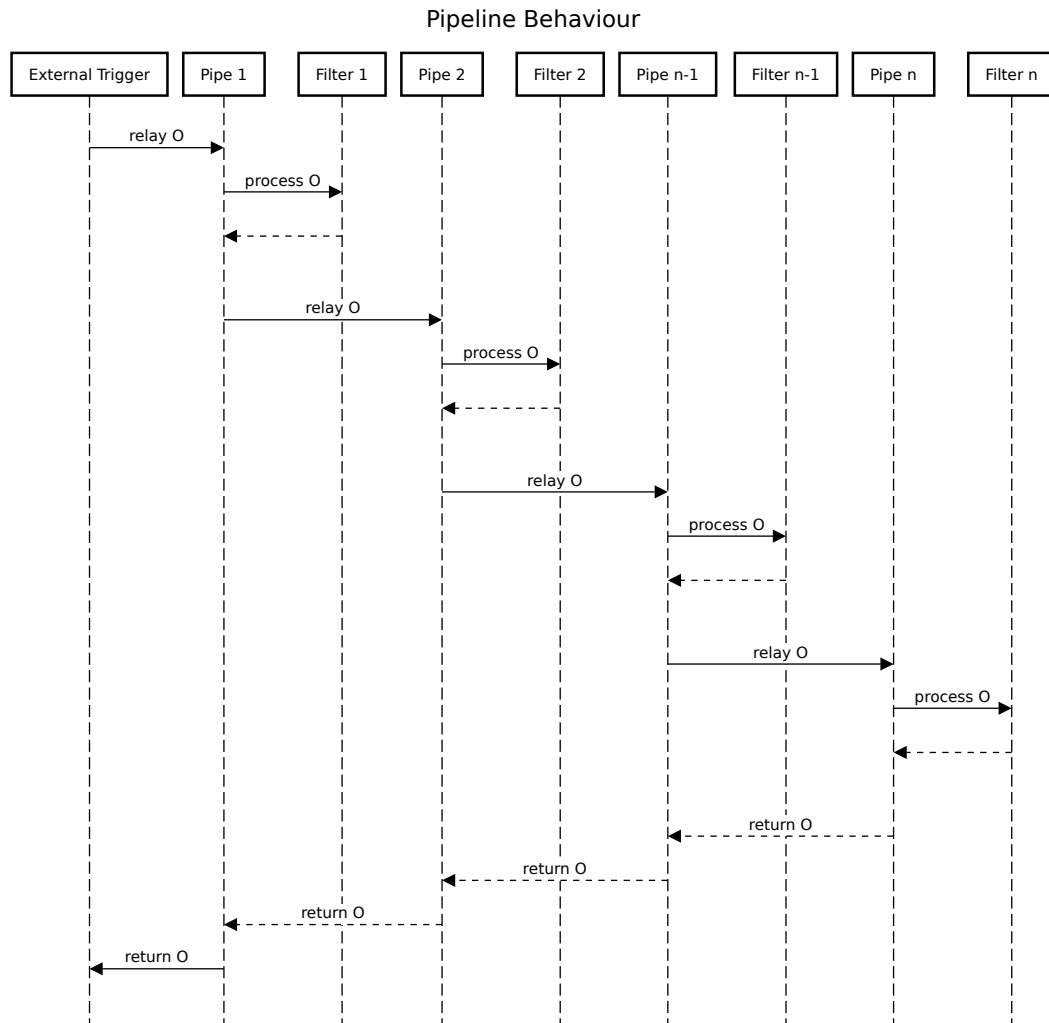
### 3.1. Design Patterns

The following sub-sections introduce a set of design patterns that are of relevance to this work.

#### 3.1.1. Pipeline/Pipes and Filters Pattern

In a paper from 1996, Alencar et al. describe the pipes and filters pattern as a mechanism to process streams of data [1]. They state that the pattern features “pipes” and “filters” components: pipes relay data streams between filters while the filters process the data streams’ contents. Figure 3.1 shows an exemplaric sequence of  $n$  pipes and filters relaying and processing an object  $O$  by implementing pipes as method calls. Alencar et al. state that the pattern’s “objective is to obtain highly reusable, interchangeable and maintainable applications”.

### 3. Theoretical Background



**Figure 3.1.:** An abstract sequence diagram of the pipes and filters design pattern. The relaying mechanism used by pipes in this diagram is realized through method calls.

#### 3.1.2. Abstract Factory/Kit Pattern

Gamma et al. describe the intent of the abstract factory pattern as follows: "provide an interface for creating families of related or dependent objects without specifying their concrete classes" [13]. They propose the following components:

- "AbstractProduct": interface for products.
- "AbstractFactory": interface for creating objects that implement "AbstractProduct".
- "ConcreteProduct": classes that implement "AbstractProduct".
- "ConcreteFactory": classes that implements "AbstractFactory" and create "ConcreteProducts".

- “Client”: class that only uses the interfaces of “AbstractProducts” and “AbstractFactories”.

They conclude that there are multiple consequences to using the pattern, one being that it “isolates concrete classes”, meaning that there is a clear isolation from the Client and the ConcreteFactories.

#### 3.1.3. Publish-Subscribe/Observer Pattern

In their book “Design Patterns - Elements of Reusable Object-oriented Software”, Gamma et al. state that the observer pattern aims to "define a one-to-many dependency between objects so that when one object changes state, all its dependents are notified and updated automatically" [13]. To achieve this, they propose the following components:

- “Observer”: interface for observing objects. It defines a single *Update* method.
- “Subject”: interface for observable objects that can (un-)register “Observers” by defining *Attach* and *Detach* methods. It can notify its registered “Observers” by use of its *Notify* method which calls all of its “Observers” *Update* methods.

Gamma et al. point out a number of benefits using this pattern. The “support for broadcast communication” is a benefit of particular relevance for this thesis as it puts a focus on the simplification of the process of sending notification to multiple, interested objects. For this reason, this pattern is used in communication protocols such as Message Queuing Telemetry Transport (MQTT) (further elaborated on in section 3.3.2).


## 3.2. Computer Communication: The OSI-Model

The Open Systems Interconnection (OSI)-Model was initially proposed in the International Organization for Standardization (ISO)/International Electrotechnical Commission (IEC) standard “ISO 7498 - Information processing systems — Open Systems Interconnection — Basic Reference Model” in 1984 and revised in 1994 by the ISO/IEC standard 7498-1 [18]. It aims to formalize a unified approach to communication between diverse peer systems by defining the *network layers* shown in figure 3.2.

### 3. Theoretical Background

---

These layers are read from bottom to top, increasing in complexity and abstraction:



**Figure 3.2.:** “Seven layer reference model and peer protocols” [18] as proposed in ISO/IEC 7498.

1. Physical: Bitwise transmission of data (e.g. via copper cable or fibre glass).
2. Data Link: Aggregation of bitwise data into data frames (e.g. via Ethernet packets sent via Wireless LAN (WLAN)) and transmission of those frames to a communication destination.
3. Network: Encapsulates data frames into packets (e.g. Internet Protocol (IP) packets) and routes and relays those packets across network nodes that are identified by addresses (i.e. IP addresses).
4. Transport: Splits packets of arbitrary lengths into transmissible packets and ensures their successful transmission (when using Transmission Control Protocol (TCP)). Also, it serves as an abstraction layer for applications that operate on higher layers.
5. Session: Nowadays part of the TCP protocol, the session layer handles establishing and terminating of connections between applications.
6. Presentation: Encoding information in a format accepted by all communication peers involved (i.e. Extensible Markup Language (XML) and JavaScript Object Notation (JSON)).
7. Application: High-level application functionality that makes use of the lower layers to communicate with peers (e.g. Hypertext Transfer Protocol (HTTP)).

This concept of a stack of network layers results in a series of encapsulated messages. For example, a large HTTP response containing a binary file can be represented as follows over the various layers:

7. Application: The HTTP response itself.
6. Presentation: The text-based HTTP headers and the binary content encoded as raw bytes.
5. (*Session: Part of the TCP protocol.*)
4. Transport: Multiple TCP packets with a binary payload and header information about the source and destination ports.
3. Network: Multiple IP packets with a binary payload and header information about the source and destination addresses.
2. Data Link: Multiple Ethernet frames with a binary payload and header information about the source and destination peer's Media Access Control (MAC) addresses.
1. (*Physical: Stream of individual bits that make up the Ethernet frames.*)

## 3.3. Internet of Things

### 3.3.1. Fields of Use

In their paper, Perera et al. categorized IoT applications into several classes [34]:

- A Smart Wearable: smart products that can be worn on different body parts or clothing.
- B Smart Home: connected applications installed and/or used in home environments.
- C Smart City: connected applications for large-scale use in cities that support logistic challenges such as traffic control and resource management.
- D Smart Environment: applications that provide monitoring capabilities for environmental metrics such as air quality and water quality.
- E Smart Enterprise: applications used in commercial and industrial environments to address challenging tasks such as logistics, transportation and energy management.

#### 3.3.2. Common Protocols

Building up on pre-existing network infrastructure and in order to meet requirements specific to individual fields of use and use-case scenarios, the landscape of IoT attends with a great variety of *communication protocols* (further used to refer to both transport and application protocols). This section will provide a brief overview of the working principles, use cases and history of some protocols commonly used in IoT and IIoT applications today.

##### ***Hypertext Transfer Protocol (HTTP)***

Initially conceived by Berners-Lee et al. at the European Organization for Nuclear Research (CERN) in 1991, the HTTP protocol is an application layer protocol that defines *requests* to resources made by clients and *responses* to said requests replied by servers [43]. According to RFC1945, HTTP requests consist of [28]:

- A request line including the HTTP verb (e.g. *GET* or *POST*), the requested resource and the HTTP version used (e.g. *HTTP/1.0*). The verb can be used to indicate what kind of action is requested (i.e. a *GET* request should retrieve contents while a *POST* request could be used to create new content). The request line is terminated by a set of Carriage Return (CR) Line Feed (LF) characters.
- A set of request header fields delimited by a set of CR LF characters where headers are encoded in the format *Key : Value*.
- An empty line indicates the end of the header fields.
- Optionally, requests can contain a message body. Its encoding is dependent on the value of the *Content-Type* header. If present, the length of the message body is specified in the *Content-Length* header.

The structure of HTTP responses is similar to HTTP requests [28]:

- A status line including the HTTP version and status code (e.g. 200 meaning “OK”, indicating a successful response). The status line is terminated by a set of CR LF characters.
- A set of response header fields, encoded just like the above-said request headers.

- An empty line.
- Optional message body. Like request message bodies, the encoding of the message body depends on the value of the *Content – Type* header.

Figure 3.3 shows a HTTP requests to the site “httpbin.org”. By definition, HTTP sends data in clear-text. Thus, HTTP communication can be intercepted and manipulated by malicious actors. Hypertext Transfer Protocol Secure (HTTPS) was introduced to solve this issue by sending HTTP requests and responses through Transport Layer Security (TLS) encrypted connections [37].

```

1 mo@ubuntu-vm:~$ curl -v https://httpbin.org
2 ...
3 > GET / HTTP/2
4 > Host: httpbin.org
5 > user-agent: curl/7.68.0
6 > accept: */*
7 >
8 * Connection state changed (MAX_CONCURRENT_STREAMS == 128)!
9 < HTTP/2 200
10 < date: Sat, 22 May 2021 19:35:20 GMT
11 < content-type: text/html; charset=utf-8
12 < content-length: 9593
13 < server: gunicorn/19.9.0
14 < access-control-allow-origin: *
15 < access-control-allow-credentials: true
16 <
17 <!DOCTYPE html>
18 <html lang="en">
19
20 <head>
21     <meta charset="UTF-8">
22     <title>httpbin.org</title>
23 ...

```

**Figure 3.3.:** A truncated HTTP request (indicated by “>”) to “httpbin.org” using the utility “curl” and the truncated response received from the remote server (indicated by “<”).

### WebSocket (WS)

The WebSocket protocol is a bidirectional application protocol used on top of TCP connections proposed in RFC6455 [22]. As stated in the RFC, it was designed to enable web applications to communicate via a bidirectional channel. In order to reuse existing infrastructure, WS was designed to upgrade existing TCP connections for HTTP and replace the text-based HTTP communication with a binary frame-based protocol. The WS protocol defines header fields for various purposes. One of those fields is the *reservedbits* field that was intended to be used with extensions

### 3. Theoretical Background

---

to the protocol. One such extension is the Per-Message Compression Extension (PMCE) that implements compression for frames [47].

#### ***Message Queuing Telemetry Transport (MQTT)***

Proposed as an Organization for the Advancement of Structured Information Standards (OASIS) standard, MQTT is described as a protocol that requires a client/server architecture and employs the publish/subscribe pattern for message queuing and routing [14]. MQTT clients can publish and subscribe to *topics*<sup>1</sup>. MQTT servers (called “brokers”) do not transform messages or process them in any meaningful way but only implement routing and broadcasting functionalities. Thus, contrary to other client/server based protocols, the central unit is not intended to process the transmitted messages. Similarly to HTTP, MQTT by itself does not implement security features and connection credential are sent in plain text. To improve on this, TLS can be used as an underlying transport protocol to encrypt MQTT messages. In order to be used in web applications, MQTT can be tunnelled through WS communication.

#### ***Industrial Protocols***

In industrial automation environments, other additional protocols are used. These include protocols that were initially conceived as serial protocols that communicated via bus topologies (e.g. Modbus and Profibus) and were over time upgraded to be operated in IP-based networks (e.g. Modbus TCP and Profinet) as well as modern protocols that employ security features (such as OPC Unified Architecture (OPC U/A)). As these protocols are of little relevance to this thesis, they are not discussed in further detail.

### **3.4. Information Security**

The following sections discuss properties, methodologies and tools used in information security.

---

<sup>1</sup>Topics are arbitrary strings. MQTT defines special characters that can be used as wildcards.



#### 3.4.1. The CIA Triad

In information security, there are three properties to information systems defined in ISO/IEC 27000:

- Confidentiality: “Property that information is not made available or disclosed to unauthorized individuals, entities, or processes”. [17]
- Integrity: “Property of accuracy and completeness”. [17]
- Availability: “Property of being accessible and usable on demand by an authorized entity”. [17]

These properties are commonly referred to as the “CIA triad” of “CIA definitions”. However, there are shortcomings to these definitions that are discussed in literature: for instance, by presenting a set of examples, Lundgren et al. point out that the CIA definitions are not suitable to appropriately cover all security requirements [21].

#### 3.4.2. Methodology

Information security involves a variety of techniques and methodologies. The ones relevant to this work are discussed in the following sections.

##### ***Penetration Testing***

The SANS Institute defines penetration testing as “the authorized, scheduled and systematic process of using known vulnerabilities in an attempt to perform an intrusion into host, network or application resources” [9].

There are various guidelines for penetration tests targetting different types of applications: The Open Web Application Security Project (OWASP) Application Security Verification Standard (ASVS), Mobile Application Security Verification Standard (MSVS) and IoT Security Verification Standard (ISVS) propose open standards for conducting web application, mobile application and IoT penetration tests, respectively [30][32][31].

The OWASP Penetration Testing Execution Standard (PTES) divides the penetration testing process into seven phases [33]:

1. Pre-engagement Interactions: Preparation of the penetration test and definition of the scope of the test (e.g. networks and duration).

### 3. Theoretical Background

---

2. Intelligence Gathering: Collecting information about the target systems (e.g. making use of Open Source Intelligence (OSINT) methods and tools) to reduce the risk of being detected by security systems.
3. Threat Modelling: Inferring potential systems in place, risks associated with the operation of those and possible mitigations thereof that may indicate potential attack vectors<sup>2</sup>.
4. Vulnerability Analysis: Identification and verification of vulnerabilities.
5. Exploitation: Active abuse of found vulnerabilities to achieve goals such as privilege escalation or pivoting<sup>3</sup>.
6. Post Exploitation: Establishing persistent access to compromised systems and exfiltration of data (such as evidence of compromise).
7. Reporting: Producing a report that documents found vulnerabilities and exploits used against the system. This document serves the contractor as a basis for decisions to make about the further development of the system.

#### ***Red-Teaming***

Oakley describes (cyber) red-teaming as a technique “to simulate [an] attack against an organization to test information systems and their related facilities” [29, p.1-14].

#### ***Fuzzing***

Sutton et al. define fuzzing as “a method for discovering faults in software by providing unexpected input and monitoring for exceptions” indicating “undefined or insecure behaviour” [42, p.22].

#### ***Man-In-The-Middle Attacks***

The Australian Cyber Security Centre (ACSC) defines Man-in-The-Middle (MITM) as “a form of malicious activity where the attacker secretly accesses, relays and

---

<sup>2</sup>“The attack vector is the way in which an attack reaches its target.” [15]

<sup>3</sup>“Pivoting is the exclusive method of using an instance also known by ‘foothold’ to be able to ‘move’ from place to place inside the compromised network.” [40]

possibly alters the communication between two parties who believe they are communicating directly with each other” [35].

#### 3.4.3. Tools

There are many tools used in information security. They vary greatly in their features, fields of use and maturity. The following paragraphs describe tools relevant to this thesis and the fields of use it touches.

**Wireshark** First released in 1998, *Wireshark* is a cross-platform and open-source tool used for network analysis, including *network sniffing*<sup>4</sup> [44]. It is written mainly in C, consists of more than 3,600,000 lines of C code<sup>5</sup> and features a Graphical User Interface (GUI). Although it is described as a network protocol analyser, it also supports sniffing of Universal Serial Bus (USB) packets. It implements a wide array of *dissectors* for various protocols and allows detailed examination of network packets (as shown in figure 3.4).

#### ***Protocol Specific MITM Tools***

The following tools are MITM that support specific protocols only:


**Burp Suite** Developed and distributed by “PortSwigger” as a commercial product, *Burp Suite* is a tool specialized for web-application testing [36]. It can be used as a MITM for HTTP communication by configuring the operating system or browser to use its internal HTTP server as a proxy. While it implements basic support for WS, it is mainly used for HTTP (and nowadays HTTPS) and lacks support for other protocols. Aside from its internal proxy server, it also provides specialized features such as the “Repeater” which is used to send forged HTTP requests. The freely available “Community Edition” (shown in figure 3.5) allows use of most of the tool’s features.

---

<sup>4</sup>“Packet sniffing is a method of tapping each packet as it flows across the network.”[2]

<sup>5</sup>This number was returned by the *cloc* utility run on commit *c73ab16b* from 23rd May 2021 of Wireshark’s GitLab source-code repository [45].

### 3. Theoretical Background



**Figure 3.4.:** Screenshot of Wireshark being executed and dissecting a HTTP *GET* request to the site “httpbin.org”. The display-filter “tcp.port == 80” shows only packets sent to or from port 80 (e.g. HTTP communication).

**mitmproxy** “A” multi-purpose MITM tool that implements (de-) serialization of TCP packets, TLS, HTTP, and WS communication [24]. It is written in Python and features a web-interface and a Python application programming interface (API) for extending mitmproxy.

**mProxy** A MQTT proxy written in Go that by itself does not feature any interactive interfaces for penetration testers and lacks support for MQTT over TLS [26]. It is, however, scriptable and enables penetration testers to execute scripts on individual MQTT messages.

**IOXY** Similarly to mProxy, “IOXY” is a MQTT proxy written in Go. It improves on mProxy’s shortcomings and implements a web interface and supports MQTT over TLS [16].

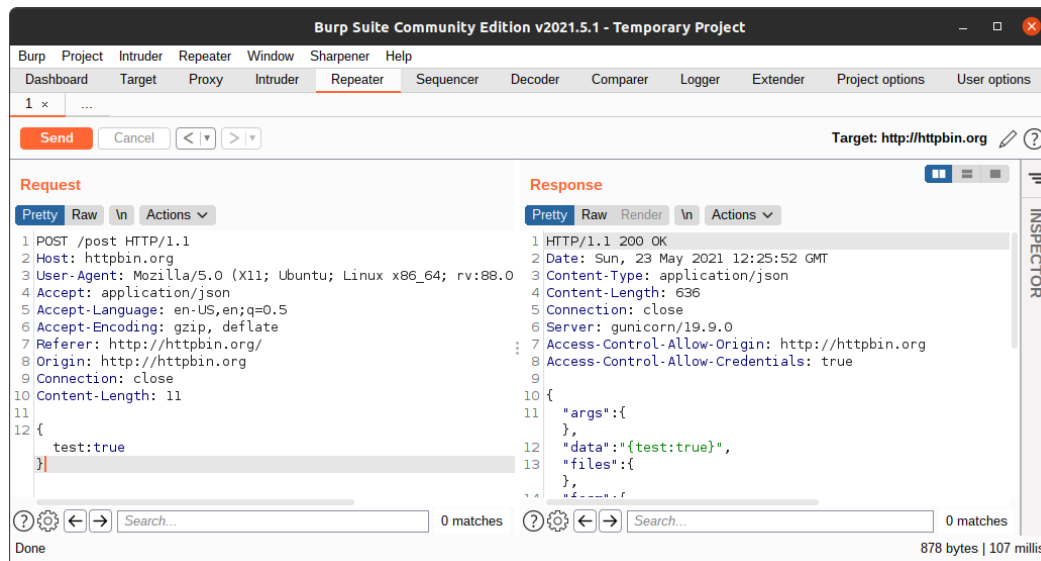


Figure 3.5.: Screenshot of Burp Suite being used to send forged HTTP requests to the site “httpbin.org”.

### Generic MITMs Tools

The following tools are generic MITM tools that support a wide range of network protocols:

**ettercap** While *ettercap* was initially developed as a network sniffer for switched Local Area Network (LAN), it was gradually extended to implement a set of MITM attacks such as Address Resolution Protocol (ARP) spoofing and *packet filtering* which allowed modifying intercepted communication [12]. Penetration testers can write custom filters in a scripting language to implement their own packet filtering logic. It is written in C and implements network protocols of layers 1 to 4 of the OSI model. Thus, it does not implement application protocols.

**bettercap** Similar to *ettercap*, *bettercap* implements network sniffing and other features used for network analysis and discovery. However, contrary to *ettercap*, it aims to support a wider range of transport technologies and is described by its authors as “*the Swiss Army knife for WiFi, Bluetooth Low Energy, wireless HID hijacking and IPv4 and IPv6 network reconnaissance and MITM attacks*” [8]. It is written in Go and features a web-interface for configuration, control and monitoring.

### 3. Theoretical Background

---

**Scapy** Its authors describe “scapy” as a “a powerful interactive packet manipulation program” [39]. It is written in Python and can be used to craft and parse packets of various protocols.

**MITMf** Built on top of scapy, “MITMf” implements a set of attacks against network devices [23]. However, it is no longer updated and its authors suggest using bettercap instead.

## **Chapter 4**

### **Procedure**

The following chapters document the work on the design concepts and implementations of the modular proxy application in chronological order. However, the common thread running through these chapters can be hard to miss which is why this chapter briefly elaborates on the procedure of and steps taken in this work.

To get an understanding of the problem space, a first prototype was built. The challenges that were encountered during implementation resulted in a set of questions towards experts. These very specific questions and a set of generic questions regarding the context of assessing the security of IoT applications were what became an interview guideline. This guideline was used to conduct interviews with a set of experts and answer the aforementioned questions.

These answers and the insights gained from designing and implementing the first prototype were then considered during two consecutive iterations of developing a new design concept. The resulting design concept was then used to implement another prototype.





## **Chapter 5**


# **Understanding the Problem Space**

In order to provide a satisfying solution to the problem at hand, the problem itself and the environment it occurs in must be researched. This chapter aims to explore and examine the problem space, resulting in a set of artefacts (namely a scenario and a set of requirements) that aid in understanding the context and designing an appropriate solution. First, a prototypical network proxy is designed and implemented in section 5.1 to get an understanding of the problems and challenges involved in designing, implementing and using such software. Based on these experiences, interviews with experts in penetration testing are conducted and evaluated in section 5.2 to get a proper understanding of their everyday work and accompanying problems. Lastly, existing software that aims to intercept communication for various scenarios and technologies is compared in section 5.3 and their usefulness for the problem-specific scenarios is assessed.

### **5.1. Prototypical Implementation**

The prototype was designed to be used in three different scenarios, each taking place in a different context. The goal of this section was to implement a prototype that could be used as a proxy to intercept communication between an IoT device and its cloud service as shown in figure 5.1. It was developed incrementally so individual components could be derived from requirements, designed, implemented and evaluated in fixed sprints.

## 5. Understanding the Problem Space



**Figure 5.1.:** Installing a MITM proxy to intercept network communication for penetration testing.


### 5.1.1. Example Scenarios

The following scenarios describe hypothetical configurations of IoT/IIoT devices that should be tested with the prototype and vary in both technical and logical complexity as well as in closeness to reality:

#### **Scenario #1: Legacy Industrial Control System (ICS) Application**

In this IIoT scenario, a Human-Machine Interface (HMI) (e.g. *Siemens KTP400 Basic*) sends commands to and receives data from a Programmable Logic Controller (PLC) (e.g. *Siemens S7-1200*) using Modbus TCP (depicted in figure 5.2). The PLC continually counts up a value up to 100 and begins anew at zero while the HMI displays the current value and provides a button that, upon being pressed by a user, resets the current value to zero.

In this scenario, attackers could perform a variety of attacks on the system by intercepting and manipulating network traffic, for example:



**Figure 5.2.:** The simplistic architecture of the first scenario.

- By dropping messages sent from the PLC to the HMI, the application may appear unresponsive as new data is not displayed on the HMI. In production environments, this could lead to dangerous situations as sensor readings that indicate harmful environmental conditions would not be presented to super-

vising personnel (e.g. sensor readings indicating dangerously high pressures of a gas tank).

- When dropping messages sent from the HMI to the PLC, control commands can be suppressed. This attack can result in catastrophic situations when emergency shutdowns issued by supervising personnel are not registered by the affected machines.

Although this scenario involves a rather simple process, it depicts a realistic communication configuration. It focusses on the use of a legacy transport protocol. Due to the rather simple structure of the Modbus TCP protocol, intercepting and manipulating communication is expected to be trivial.

### Scenario #2: IoT Cloud Application

As shown in figure 5.3, this IoT smart home scenario utilizes two local IoT devices that are integrated into a cloud environment such as the Amazon Web Services (AWS) IoT platform: a thermometer and an Air Conditioner (A/C) unit. Both devices connect to the cloud platform, authorize themselves at a Representational State Transfer (REST) interface via HTTP and upgrade their HTTP connection to WS streams upon successful authorization. They eventually communicate to a remote MQTT broker by tunnelling MQTT packets over the WS stream. At this stage, the thermometer publishes temperature readings to an MQTT topic while the A/C unit subscribes to the same topic and adjusts its operation depending on the incoming temperature readings.

If not securely configured, this distributed communication setup introduces a set of possible attacks that could be performed when attackers *impersonated* client-devices or the remote server:




**Figure 5.3.:** The central MQTT broker relays messages to subscribed clients. In this case, the A/C unit receives messages published to a topic it subscribed to.

- Impersonating the thermometer, attackers could send incorrect temperature data and effectively control the A/C unit. When sending low temperature readings while the environment temperature is high, the A/C unit would stop running. Conversely, when high temperature readings are sent while the en-

## 5. Understanding the Problem Space

vironment temperature is low, the A/C unit would run, and thus further cool down the environment.

- Attackers that impersonate the remote server could drop or manipulate incoming publish packets, thus altering whether and/or what information is relayed to other connected devices. For example, temperature readings that indicate a high environment temperature that would lead to the A/C unit to be powered up could be rewritten in such a way that the transmitted temperature value is considered to indicate a low environment temperature, thus preventing the A/C unit from running automatically.



**Figure 5.4.:** State machine of AWS IoT communication


This scenario makes use of three communication protocols, uses these protocols dependent on the state of authentication and even tunnels one protocol through another one. Therefore the proxy application has to implement a state-machine (as seen in figure 5.4) and testing communication in this scenario is expected to be more complex than the first one. Also, since this scenario makes use of the AWS IoT infrastructure, special care must be taken to ensure that authentication is properly relayed to the cloud servers.

### **Scenario #3: Water Treatment Plant**

Similar to scenario #2, this scenario makes use of MQTT for transporting messages. However, the scenario takes place in an ICS context of critical infrastructure.

As shown in figure 5.5, there are multiple steps involved in treating water for drinking. The scenario represents these steps as separate stations (“source”, “flocculants”, “sedimentation”, “filtration”, “disinfection” and “storage”) that act as MQTT clients. Each station receives water into an input tank, processes water from its input tank in a specified rate and flushes processed water into an output tank. Similar to how threads can suffer from starvation in a multithreading environment, these

<sup>1</sup>[https://en.wikipedia.org/wiki/File:Illustration\\_of\\_a\\_typical\\_drinking\\_water\\_treatment\\_process.png](https://en.wikipedia.org/wiki/File:Illustration_of_a_typical_drinking_water_treatment_process.png)



**Figure 5.5.:** Illustration of a typical drinking water treatment process. (by the CK-12 Foundation)<sup>1</sup>

stations can either “run dry” when their input tank is empty or overflow when either tank is filled beyond their capacity. In this scenario, stations will only process water from their input tanks if their output tank provides sufficient available capacity.

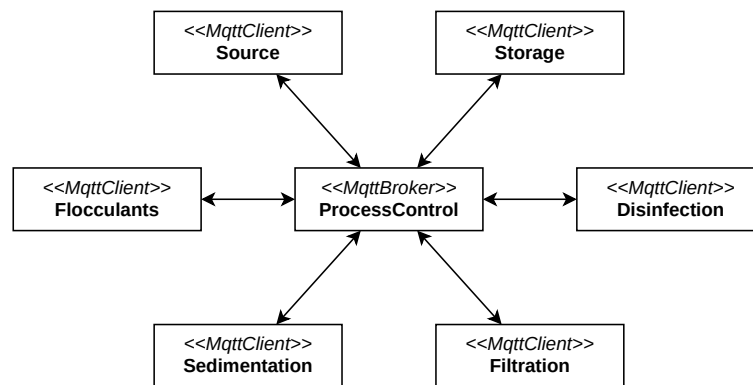
Similar to scenario #2, attackers could perform a series of attacks in this scenario:

- Without intercepting any communication, attackers could cyclically overwrite water levels of either input or output tanks to stop stations and bring processing to a halt. For example, when attackers overwrite the “storage” station’s input tank level to indicate exhausted capacities, the “disinfection” station would fill its output tank and eventually stop processing water. This would cause the “disinfection” station’s input tank to fill up and lead to the “filtration” station’s output tank to fill up. Ultimately, the water treatment plant would halt.
- When any station publishes data about its tanks’ levels indicating full or empty capacities, attackers could intercept those messages and change them to either indicate the opposite (change tank levels indicating full capacities to levels indicating empty capacity) or some arbitrary level information. This could lead to either pumping water from empty tanks, potentially damaging

## 5. Understanding the Problem Space

the pumps, or to overfilling tanks, leading to leaking excess water and potentially damaging further equipment.

- Attackers that intercept messages between the stations and the broker can collect and analyse them and try to draw conclusions about the use and activity of the system. This may allow attackers to identify patterns that show when the plant is operating at high capacities, maximizing the effect of attacks against the plant.



**Figure 5.6.:** MQTT clients for each station of the water treatment plant publish messages about their tanks' levels to individual topics and subscribe to their succeeding station's tank levels in order to simulate flushing water into their input tanks.

This scenario greatly simplifies drinking water treatment by reducing the process to the producer-consumer problem known from multithreading. A more realistic representation of drinking water treatment plants would take further details of the operations of such plants into account.

This scenario involves only MQTT as a transport protocol but, as can be seen in figure 5.6, it requires six MQTT clients to run simultaneously.

### **Derived Use-Cases**

Summarizing the scenarios detailed above, a number of high-level use-cases can be derived from them (shown in 5.7). The actors are the *attacker* that intends to interact with the system in a potentially malicious way and the *server* and *device* that use the system for transportation of messages. The following use-cases are found in the aforementioned scenarios:

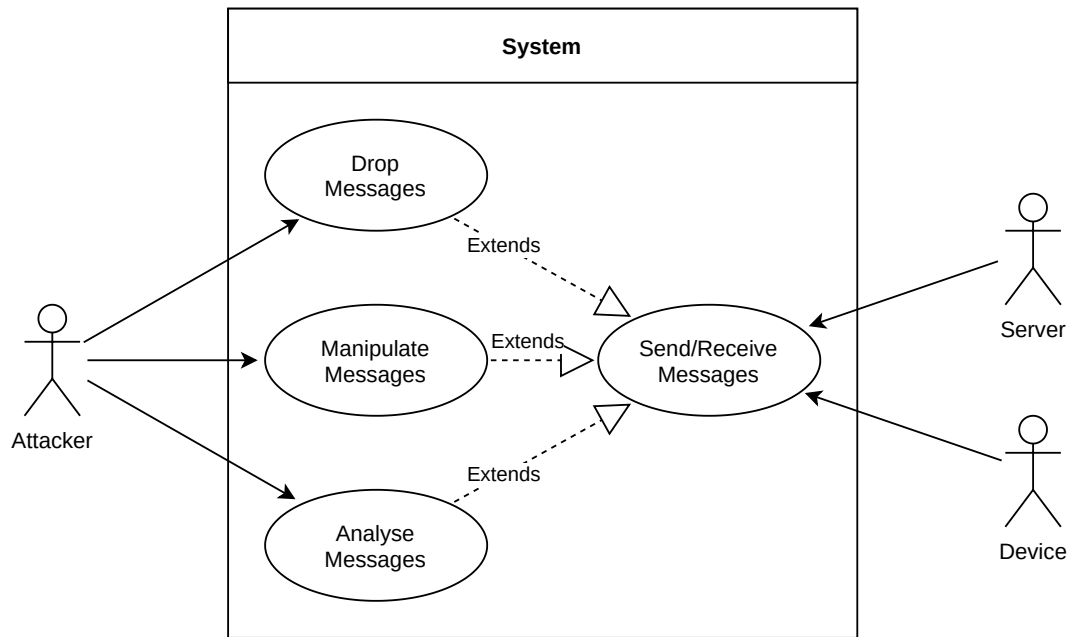


Figure 5.7.: High-level use-cases of a proxy in a generic IoT/ICS environment.

- **Send/Receive Messages:** The server and device send and receive messages to communicate with each other. This interaction does not require interaction with the attacker.
- **Drop Messages:** The attacker discards incoming or outgoing messages by not relaying them to the intended target. This can cause loss of control in the application that this communication takes place in.
- **Manipulate Messages:** Incoming or outgoing messages can be changed by an attacker, altering various properties such as Quality of Service (QoS) (for MQTT messages), host (for HTTP requests) or the payload of a message (e.g. the content of an HTTP response).
- **Analyse Messages:** Attackers can collect and analyse messages passively without altering them, allowing them to deduce information about the behaviour of the affected system and potentially its user(s).

### 5.1.2. Requirements

To be able to operate in all of the aforementioned scenarios, the prototype had to implement a set of functional requirements:

## 5. Understanding the Problem Space

---

|   |                  |
|---|------------------|
| <b>F1</b>   | <b>Protocols</b> |
| The software must implement parsing/crafting messages/packets of the communication protocols: HTTP <sup>2</sup> , WS, MQTT and Modbus TCP.                                  |                  |
| <b>Fit criterion:</b> The software must implement and support the HTTP, WS and MQTT protocols so that messages of those protocols can be further processed by the software. |                  |

|   |                       |
|---|-----------------------|
| <b>F2</b>   | <b>Network Stacks</b> |
| The software must be able to parse protocols that are tunnelled through other protocols (“ <i>stacked</i> ”). It must provide an interface to the user where they can specify which communication protocols are used and whether and how they are stacked (further referred to as <i>network stack</i> ). |                       |
| <b>Fit criterion:</b> The software processes a configuration file that lets users specify which protocols are used and whether/how they are stacked.  |                       |

|  |                       |
|--|-----------------------|
| <b>F3</b>  | <b>State-Machines</b> |
| The software must be able to switch network stacks and scripts for processing dependent on configurable <i>states</i> and <i>transitions</i> between them. It must provide an interface for the user to specify when to switch to using another network stack, represented using Finite-State Machines (FSMs) and rule sets for transmission between states. |                       |
| <b>Fit criterion:</b> The software processes a configuration file that lets users specify when to switch between network stacks.   |                       |

|   |                    |
|---|--------------------|
| <b>F4</b>   | <b>Integration</b> |
| The software shall provide interfaces for integration of third-party software.  |                    |
| <b>Fit criterion:</b> The software implements interfaces that allow sending packets to other applications such as “Burp Suite”. |                    |

|  |                  |
|--|------------------|
| <b>F5</b>  | <b>Scripting</b> |
| The software shall provide scripting capabilities for automated manipulation and discarding of messages. |                  |
| <b>Fit criterion:</b> Users can define scripts that are executed on messages.                            |                  |

<sup>2</sup>HTTPS was deemed relevant as the prototype as of an academic nature and the addition of Secure Socket Layer (SSL) introduced further complexity.



|  |                |
|--|----------------|
| <b>F6</b>  | <b>Logging</b> |
| The software shall provide means for collecting and saving messages for future analysis. |                |
| <b>Fit criterion:</b> The software saves messages to a MySQL database.                   |                |

The following non-functional requirements were defined:

|   |                               |
|---|-------------------------------|
| <b>N1</b>   | <b>Platform Compatibility</b> |
| In order to support a broad spectrum of target platforms, the software shall be implemented platform-independently. |                               |

|   |                    |
|---|--------------------|
| <b>N2</b>   | <b>Reusability</b> |
| The software shall be reusable so it can be used in future tests that may feature new configurations of network stacks. |                    |

|  |                    |
|--|--------------------|
| <b>N3</b>  | <b>Open Source</b> |
| The software shall be available as open source software so programmers and members of the IT community may contribute to improving it. |                    |


Due to this implementation serving as a prototype and being of an academic nature, no specific constraints were defined. It was to be developed strictly ignoring aspects of usability and stability as it should not be used in production environments but in laboratories exclusively.

### 5.1.3. Design

The prototype was designed to be fit for use in the second scenario as, regarding network communication, it was more complex than the other ones. Specifically, the second scenario demanded the implementation of a network stack and a state machine to switch between states. Parsing protocols that were tunnelled through other protocols appeared to be a potentially challenging requirement which is why the focus on the design and implementation of this prototype was on the underlying management, processing and relaying of messages. In order to tackle it, a variation of the *pipeline* (sometimes referred to as *pipes and filters*) design pattern was used (as shown in figure 5.8). It was designed to be used as follows:

## 5. Understanding the Problem Space

*Messages* originate from a listener, for example messages with raw byte payloads are received from a TCP socket. These messages are sent to an initial *pipe* to be processed *down*.



**Figure 5.8.:** The variation of the “pipes and filters” design pattern used in the prototype.

Pipes are bi-directional routers that represent processing-steps of pipelines and perform the following actions on messages that are processed through a pipeline:


1. Pipes use optional *encoders* to disassemble/de-serialize messages when processing them *down* the pipeline and re-assemble/serialize them when they process messages *up* the pipeline.
2. Pipes can use *filters* to perform operations on messages such as replacing header values or manipulating payloads.
3. They forward messages to the next pipe in its pipeline when processing messages down or to the previous pipe when processing messages back up.

There are extensions to basic pipes such as:

- *EndPipes* are appended to the end of a pipeline and reverse the message processing direction so messages that were processed down are sent back up the pipeline to be processed up.
- *ProcessingPipes* mandate encoders and filters to be used. These pipes are used to indicate that messages are not only routed but also processed and encoded or decoded.
- *IntegrationPipes* allow integration of other software into the pipeline. For example, penetration testing software such as Burp Suite could be integrated.

An exemplaric configuration of the pipeline design pattern envisioned for this prototype for use in the AWS IoT scenario is shown in figures 5.9 and 5.10. These

## 5. Understanding the Problem Space



**Figure 5.9.:** AWS IoT Scenario - State 1: HTTP Server

diagrams visualize how messages are processed *down* and back *up*.

Figure 5.9 shows the first state of the AWS IoT scenario that processes HTTP communication. It features a TCP server socket that accepts incoming connection requests from an IoT device and a client socket that is connected to the AWS cloud.

Since the communication to the AWS cloud is TLS-encrypted, it is first decrypted by a filter and then processed by a HTTP filter. Then, the parsed HTTP requests and responses are sent to external tools (e.g. Burp Suite). Once the end of the pipeline is reached, the messages are sent back up the pipeline, being encoded back into a form usable for the IoT device or cloud server.

Once the prototype detects that the state must be changed to processing MQTT over WS communication, a different network stack is initialized and used, as shown in figure 5.10. In this state, TLS-encryption is decrypted and passed into a WS filter that (de-)serializes WS packets. The payload of data frames is then forwarded to an MQTT layer. In this specific configuration shown in figure 5.10, the payload of MQTT messages is (de-)serialized as JSON before being sent to external tools by the integration pipe.

### 5.1.4. Testing

To test the prototype, a simple testbed was designed and implemented to realize scenario #3 (discussed in section 5.1.1). It consisted of two Debian 10 machines that acted as a MQTT broker and clients and a Kali Linux machine that ran the prototype and provided tools such as Wireshark that could be used for debugging and monitoring network traffic. All machines were connected to a single network (shown in figure 5.11) and were assigned static IP addresses. While this setup allowed for more sophisticated MITM mechanisms such as ARP spoofing, the decision was made to configure the MQTT clients to directly connect to the *kali* machine to reduce complexity and accelerate and simplify testing. Separate machines were used for the MQTT broker and clients so that actual MITM attacks could be performed if the need to arose. Also, running the broker on a separate machine simplified debugging as network traffic could be attributed to broker or clients easier by examining the packets' source and destination IPs.


The MQTT broker software used on the *broker* machine was Eclipse Mosquitto<sup>3</sup> 1.5.7 and had the WS transport enabled to allow clients to connect via WS. The MQTT clients running on the *client* machine were implemented in Python using the Eclipse Paho library for Python (paho-mqtt<sup>4</sup>).

---

<sup>3</sup><https://mosquitto.org/>

<sup>4</sup><https://pypi.org/project/paho-mqtt/>


## 5. Understanding the Problem Space




**Figure 5.10.:** AWS IoT Scenario - State 2: MQTT via WS

The water treatment scenario required water treatment stations to be simulated individually as separate MQTT clients, which was done by representing them as “Pro-

## 5. Understanding the Problem Space




**Figure 5.11.:** A network diagram of the testbed that was used for testing the prototype.



**Figure 5.12.:** The “ProcessingUnit” data-structures represent individual stations of the simplified water treatment plant.


cessingUnits” in the Python implementation of the testbed. As can be seen in figure 5.12, ProcessingUnits held individual *MqttClient* instances running in separate threads, were subscribed to the topics of relevant other units such as their direct predecessors and successors and were capable of publishing their current state. Their *process* method would be called cyclically and allow for units to calculate their intake, throughput and output.

These units were then “chained” up (shown in figure 5.13) in the order in which they were presented in the scenario by specifying their direct predecessor and successor units: potentially contaminated water would be pumped out of the *source*, processed by a series of stations and eventually flushed into the *storage*. The *source* was an instance of the “SourceUnit” that featured a throughput calculated by a sine-wave function that used the elapsed time since program startup as input parameter. Also, in order to keep the program running infinitely without either the *source* “running



**Figure 5.13.:** Chaining of the water treatment units, originating from a water source and eventually leading to a storage at the end of the processing pipeline.

dry” due to its input tank emptying or the *storage* overfilling, the *storage*’s output was programmed to feed back into the *source*’s input tank (as can be seen in figure 5.14). While this was not a realistic approach, it kept the program’s design simple and allowed for continuous testing and did not impact the MQTT communication.




**Figure 5.14.:** Screenshot of the application “MQTT Explorer” that was used to inspect and visualize the state of the water treatment plant. The left graph shows how the *source*’s input tank steadily emptied until it was filled by the *storage*’s output tank. The right graph shows how the *flocculant* unit’s input tank slowly filled up.

### 5.1.5. Implementation

The prototype was partially implemented over the course of 8 weekly sprints after which work on the prototype was halted due to time constraints. It was written in TypeScript due to the language’s flexibility and static typing. It allowed to precisely specify interfaces and its runtime (NodeJS) would allow it to make use of asynchronous programming, which would benefit this prototype as this simplifies working with asynchronous tasks (such as networking). The rough design worked out in section 5.1.3 was specified in greater detail so individual classes could be derived and implemented.

**Pipes and Filters** As shown in figure 5.15 the pipeline design pattern was altered in such a way that the basic *IPipe* interface was implemented by the *BasePipe* class that held a reference to a single *IFilter* but lacked a reference to an *IEncoder*.

## 5. Understanding the Problem Space



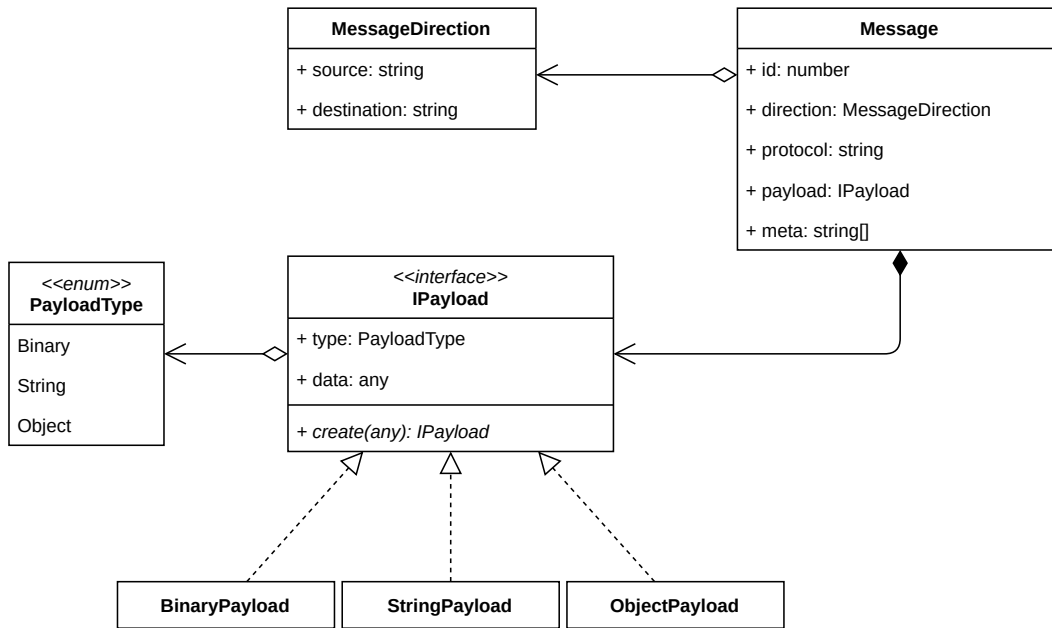
**Figure 5.15.:** The classes and interfaces used to implement pipelines in the TypeScript prototype.

Filters would hold a reference to encoders because encoders were used directly by filters for (de-)serialization prior to any other processing (such as executing scripts). Thus, encoders would not exist without filters, resulting in a composition relationship between the two. As indicated by their prefix, *IFilters* and *IEncoders* were only interfaces that set a behaviour for their specific uses: encoders would implement (de-)serialization of specific protocols while filters added logic to processing (de-)serialized packets, such as executing scripts. The prototype implemented encoders for HTTP and MQTT as well as a *BaseFilter* that did not add any logic to processing but allowed to test the encoder implementation. Also, specific *NopFilter* and *NopEncoder* (“Nop” meaning “no operation”) classes were implemented that did not implement any logic. This was used to test sending messages down and up the pipeline without processing them at all. The MQTT encoder used the “mqtt-packet”<sup>5</sup> library which offered a comparatively simple API for serializing (*generate(Packet)*) and de-serializing (*parser.parse(Buffer)*). However, lacking a library that offered a similar low-level and simple API for HTTP (de-)serialization, a custom encoder for these tasks was implemented. Due to HTTP being a comparatively simple and text-based protocol, all that needed to be done for de-serialization was parsing the HTTP headers (separated by new-lines) and, depending on whether or not the “Content-Length” header was present, extracting the HTTP body.

**Messages and Payloads** The pipeline system implemented routing and processing of messages, however this required a concept of what messages are and how

<sup>5</sup><https://github.com/mqttjs/mqtt-packet>, commit 4b6278d890e0c2fca01da62c5f9b63e05f5fd899



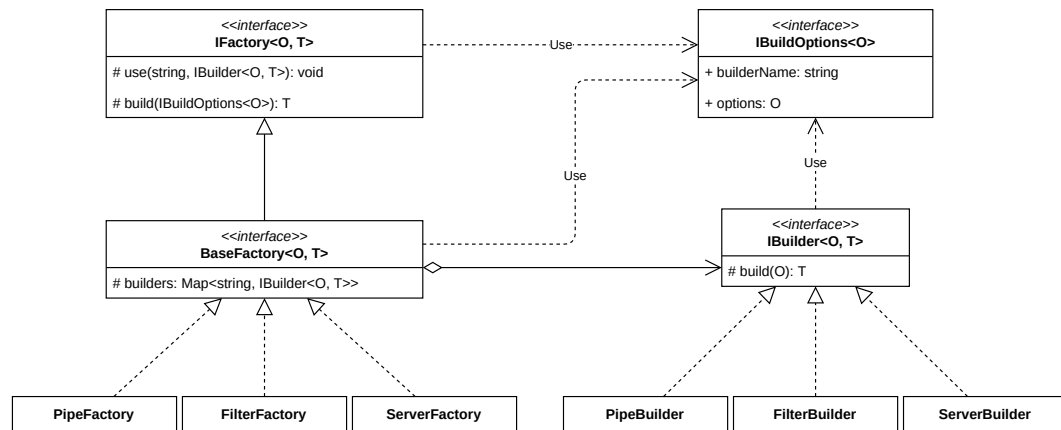


**Figure 5.16.:** The structure of messages in the prototype: messages hold information about their communication direction, protocol-specific meta data and type-dependent payloads.

they convey the information required to perform meaningful and useful operations on network communication. Figure 5.16 shows the classes and interfaces that defined the messages and payloads types. *Messages* hold basic information such as a unique identifier, the communication protocol they were sent through and metadata that was used to store header information in. The *IPayload* interface allowed for implementation and use of various payload formats such as raw binary information (e.g. MQTT message bodies), string contents (e.g. HTTP response bodies when the *Content-Type* header indicated *text* data) or JavaScript objects such as dictionaries that could hold arbitrary data for cases where there was no meaningful way to extract payloads from messages. The *MessageDirection* structure was used to relay the message to the correct socket after it was processed by the pipeline.

**Factories and Builders** The requirement “F2 Network Stacks” implied a way of initializing various objects that represent pipelines, sockets and FSMs, dependent on configuration files loaded at runtime. As shown in figure 5.17, the abstract factory design pattern was used to provide an easy way to initialize pipes, filters, encoders and sockets by providing them with metadata used for object creation. An *IFactory* interface exposed simple methods for registering *IBuilders* and building objects. The generic type parameters *O* and *T* were placeholders for type specific options (e.g. *IPipeOptions* and *IFilterOptions*) and the type of the created objects

## 5. Understanding the Problem Space



**Figure 5.17.:** A variation of the abstract factory design pattern was used to decouple the creation of objects from their usage.

(e.g. *IPipe* and *IFilter*), respectively. The options types would contain information that was used for creating individual instances, such as a pipe's name or a server socket's address to listen on. The *BaseFactory* class implemented the *IFactory* interface and held an internal hash-map that was used to register *IBuilders* by name. Lastly, the *IBuilder* interface provided a method for initializing objects with the given options, providing default values for constructor parameters. There were static instances of factories and builders of pipes, filters and servers. For instance, the global server-factory *SERVER\_FACTORY* used the global TCP server builder instance *TCP\_SERVER\_BUILDER*.

### 5.1.6. Insights Gained

The following insights were gained through the prototypical implementation. Some resulted in questions relevant for the expert interviews that were to be held:

- Due to the Maximum Transmission Unit (MTU), large messages were broken into chunks that were transferred sequentially. This required the proxy to work on streams of incoming data and reassemble messages from said chunks. While individual MQTT messages would often be short enough to be transmitted in a single TCP packet, other communication protocols such as HTTP could yield messages that were several hundred kilobytes or more in size (e.g. when downloading images). This also required the encoders to be stateful as they had to load data of incoming packets into individual buffers

until they could parse complete messages, introducing the need to initialize one pipeline per device connected to the proxy application.

- Supporting multiple client devices was non-trivial as communication between clients and servers was not necessarily connection-oriented (e.g. HTTP on the application level or User Datagram Protocol (UDP) on the transport level).

*Q1: Do penetration testers need to test multiple devices at the same time?*

- In some cases, e.g. with WS data-frames, extending a message's payload resulted in its payload being split into multiple messages. This indirectly created new messages that, depending on the exact protocol used, needed to use generated values (such as a unique identifier) or context-specific information (e.g. authentication tokens used in HTTP headers). Also, some libraries would generate those values themselves and not define ways to specify those manually.

*Q2: Do penetration testers require exact control over the implementation of protocols?*

- Manipulating messages, automatically via scripting or by hand using third-party integrations (e.g. to *Burp Suite*), could introduce latency to the communication.

*Q3: Are there strict timing requirements during penetration tests?*

- Many libraries offered high-level functions to the programmer while avoiding exposure of low-level functionalities like crafting or parsing messages. Exposing such functionalities would require dissecting and altering libraries on a source-code level.

### 5.2. Interviewing Experts for Insights

Interviews may be an efficient way to get an expert's opinion on something they are proficient in. Thus, expert interviews were conducted to let security researchers give insight into their everyday work and the challenges they face when working with IoT and IIoT applications. The information and insights gathered in these interviews were then used to verify various assumptions made during the design and development of the prototype and derive additional requirements.

### 5.2.1. Interview Guideline

An interview guideline (shown in appendix B) was created to keep focus on key points during interviews so that interviewees would not stray too far from the relevant points. The guideline also served as a checklist so the interviewer could make sure that all questions and points that should be covered initially, were in fact covered by the end of the interviews. It was composed of three sections:

**1. Experiences with IoT** The answers to these questions would give insights into what kind of applications the security researchers had worked on in the past. Answers to question *1.1.* were of particular interest as they might represent what technologies were being examined by security researchers and may be popular in today's applications.

**2. Processes in Everyday Life** This section aimed to cover questions about the processes and tasks security researchers perform during penetration tests of IoT applications in their everyday life. Ideally, answers to those questions would show the approaches taken and challenges faced during their work, uncovering potential needs and underlying motivation.

**3. The Future of IoT** This section had security researchers assess what the future of IoT may be like from their point of view. This required the interviewees to make a critical assessment of the status quo.

### 5.2.2. Conducting Interviews

Interviews were conducted with six *NVISO* employees (Patrick Eisenschmidt, Cédric Bassem, Théo Rigas, Oliver Nettinger, Pierre-Alain Mouy, Jonah Bellemans) that all had worked on security assignments on IoT or IIoT applications in the past. There is considerable variety in

- the experience they had in working on security assignments in general: all interviewees had a strong background in cyber security that reached back multiple years except one who was a working student at *NVISO Labs* (Bellemans).

- the experience they have had in working on IoT/IIoT applications: two interviewees worked on assessing IoT/IIoT applications only occasionally (Eisen-schmidt, Mouy), one was part of a car manufacturer's automotive security team in the past (Nettinger) and two were part of *NVISO Labs* and worked with smart devices on a regular basis (Bassem, Rigas).
- the focus of their everyday work: two interviewees were *NVISO* chief executives and switched to working on management tasks rather than security assessments (Nettinger, Mouy), one was a working student finishing their master's thesis with a focus on legal aspects of IoT devices (Bellemans) and the remaining three worked on security assessments in a variety of fields (Eisen-schmidt, Bassem, Rigas).

The duration of the interviews varied from 45 minutes to two hours depending on the amount and level of detail of information provided by the interviewees and the number of times that the interviewer had to ask further questions.

Due to the COVID-19 pandemic, interviews were conducted remotely over Microsoft Teams and recorded for later review and analysis. All interviews were conducted successfully, however some problems occurred: due to unstable internet connections interviews were sometimes interrupted for up to 30 seconds, low bandwidth and low microphone quality sometimes made making out specific words and phrases very hard.

### 5.2.3. Interview Analysis

The answers interviewees gave to the various questions in the interview guideline varied greatly in detail. The following paragraphs attempt to summarize the essential statements interviewees made, sorted by the sections of the interview guideline and ending with conclusions drawn from the interviews.

**1. Experiences with IoT** Asked about the technologies they encountered in their work, most interviewees stated that MQTT [11, 5, 38, 27] and HTTP [11, 5, 38, 27, 25, 6] were widely used in the smart applications they assessed. For IoT devices they found that Espressif microcontrollers such as the ESP32 and ESP8266 were used [11, 6]. Especially in cheap devices they found that custom protocols and infrastructure were used [5, 38, 6], whereas high-end devices usually used MQTT and

HTTP and worked with well-known cloud infrastructures such as AWS, Microsoft Azure or Google Cloud Platform. Most interviewees worked on Smart Home products [11, 5, 38, 25, 6] with one notable exception being Nettinger who worked on Smart Cars [27].

Usually, there were no technical constraints for the interviewees when performing security assessments. There were some non-technical constraints such as working from a black-box perspective rather than working from a white-box perspective that would allow evaluating more security aspects of a system in less time [11, 5, 38]. Depending on the client and the exact application that was to be tested, interviewees said that they made use of either mobile lab environments [5, 38] or stationary lab environments [11, 27, 25, 6]. Also, interviewees stated that they usually assessed devices and applications individually.

**2. Processes in Everyday Life** Regarding the goals of their assessments, interviewees would take on one of two approaches: The first was penetration testing [11, 5, 38], aiming to evaluate as many components of a system as they could during their assessment. The second was red-teaming [27, 25, 5, 38] which aimed to get some level of access, preferably privileged, to a device or server in order to take influence on the application's logic or exfiltrate data. The scope of their assessments was usually defined by the client and could include testing of devices, applications and firmware or performing source-code and cloud configuration reviews [11].

The high-level tasks carried out during assessments would generally be the same across assessments: first, interviewees would inspect applications passively from a black-box perspective without interacting with them. This could incorporate looking for hardware interfaces on a device [5, 38], looking for open network ports [6], reverse engineering Android applications and inspecting certain artefacts as manifest files [11] and monitoring applications' network traffic [6]. Nettinger stated that when working with cars, fuzzing was a task often carried out against bus protocol implementations because the devices implementing those protocols were often supplied by third parties and source-code was usually not available [27].

The tools used by the interviewees were mostly dependent on the technologies and protocols they worked with, such as Burp Suite for examining HTTP communication [11, 5, 38, 27, 25, 6]. However, some general tools were used for information gathering, monitoring (Wireshark) and networking (such as socat<sup>6</sup> and mitmproxy).

---

<sup>6</sup>Command-line based tool that is used for many networking tasks, such as forwarding network traffic [41].

Bassem, Rigas and Mouy stated that they would occasionally implement their own tools or scripts when they found that there either were no tools available that suited their needs or those tools would not work [5, 38, 25]. According to Rigas, tools were highly specific to custom setups and preparing them up for use could be more challenging than actually using them [38]. Bassem, Rigas and Bellemans stated that tools were often immature, making it difficult to work with them [5, 38, 6]. Speaking of their automated tests performed on smart devices, Bellemans criticized that automated tools often yielded inaccurate or incorrect results such as nmap reporting a game-server running on a smart lightbulb [6]. Also, when manipulating communication of applications, interviewees generally were not interested in manipulating metadata such as headers but focused on the messages' payloads.

**3. The Future of IoT** When asked about the current challenges the interviewees were facing working on IoT assessments, they gave very individual answers: Eisenschmidt expressed concerns about data protection and cloud environments being a rather new technology that requires engineers to securely configure them [11]. Mouy and Eisenschmidt stated that protocols and frameworks became increasingly complex and more and more devices interacted with each other, adding complexity to the security assessments [25, 11]. Also, there were a lot of custom protocols and frameworks that lacked proper tooling and were time-consuming to assess [6, 25]. When working on IoT assignments, clients often had a traditional view on the assignments and occasionally wanted the testers to perform black-box tests only although additional white-box tests would potentially help covering more components and internals of applications [6, 5, 38].

Half of the interviewees stated that cloud computing will be more important and present in the future [11, 27, 25]. They expect continued use of the comparatively old but proven HTTP [5, 38] and the well-accepted MQTT [5, 38, 27]. Regarding software development, they expect manufacturers of smart systems to involve IT security more into their development process [5, 38, 27] as well as use standardized frameworks [25, 5, 38]. However, they also stated concerns about the growing complexity of frameworks and the uncertainty of which frameworks will eventually gain wide acceptance [5, 38]. Regarding autonomous driving, Nettinger noted that current discussions about legal topics (such as the question about which party is to assume liability in case of accidents) will likely not come to an end anytime soon [27]. Concerned about security and safety aspects of future IoT applications, Belle-

## 5. Understanding the Problem Space

---

mans expressed the need for smart applications to be labelled or certified and they referred to the European cybersecurity certification framework that is being worked on by the European Union Agency for Cybersecurity (ENISA) [6].

**Conclusions** The interviews yielded a set of both very interesting and relevant insights into the interviewees' work and fields of expertise. The following insights served as a guide for further development of the proxy application:

- Smart devices often communicated via HTTP and MQTT. While the tools for security assessments with HTTP were very mature, there was a perceived lack of tools for MQTT.
- Often times, smart applications made use of proprietary protocols and infrastructure. While this was a fact the interviewees expect to be of less significance in the future, it still was of greater significance then.
- Penetration testers usually did not intend to test the protocol implementations used by applications but the contents transmitted over these protocols.
- Tools for working with specific protocols were often very immature and both installation and usage involved a lot of work.

Also, the questions raised in section 5.1.6 were answered and could be used to derive assumptions to take into account when creating a software design for a proxy application:

- Q1** The interviewees usually tested individual devices one at a time. Therefore, an assumption can be made that the software design should not aim for testing multiple devices at once.
- Q2** Interviewees stated that during tests there is a strong focus on the message payloads. They do not test protocol implementations and usually do not have to manipulate protocol headers. Thus, the software design should aim to provide penetration testers with access to message payloads rather than protocol headers.
- Q3** According to the interviewees, they did not encounter strict timing requirements such as real-time communication during their assessments. As a result, the software design should not take timing into account any may envision mechanisms that potentially introduce latency (such as queueing messages).

These insights were translated into the following, new non-functional requirements:



|   |                      |
|---|----------------------|
| <b>N4</b>   | <b>Extensibility</b> |
| To allow for future implementation of further communication protocols the software shall be implemented in a modular fashion. |                      |

|   |                   |
|---|-------------------|
| <b>N5</b>   | <b>Deployment</b> |
| To allow the proxy application to be installed and used in a repeatable and reliable way, the proxy application shall be distributed using a deployment system. |                   |

### 5.3. Analysis of Existing Software

Existing software was examined and compared to see whether they could be modified and extended to fulfil the requirements worked out in sections 5.1.2 and 5.2. Table 5.1 shows a comparison of popular software used for MITM attacks in information security. As can be seen, some tools meet exclusion criteria such as being incapable of crafting packets (e.g. Wireshark) or not being released under an Open Source Software (OSS) license (e.g. BurpSuite). Protocol specific tools such as mitmproxy, mProxy and IOXY usually do not consider support for further protocols in their software design which makes implementing new protocols to these tools a very challenging tasks. MITMf was not updated in years and lacks a clear documentation. The remaining tools (Scapy, Ettercap and bettercap) provide generic interfaces for parsing and crafting packets and executing scripts on packets. However, none of these feature FSMs for dynamically changing the network stack used for disassembling packets.

For these reasons, a decision was made to implement a new tool instead of extending an existing one.

## 5. Understanding the Problem Space

| <i>Name</i> | <i>Latest Release</i> | <i>Implemented in</i> | <i>Supported Protocols</i> | <i>R</i> | <i>W</i> | <i>D</i> | <i>OSS</i> |
|-------------|-----------------------|-----------------------|----------------------------|----------|----------|----------|------------|
| Wireshark   | 2021-04-21            | C                     | Various                    | F        | N        | N        | Y          |
| BurpSuite   | 2021-05-26            | Java                  | HTTPS, WS                  | F        | F        | F        | N          |
| Scapy       | 2021-04-19            | Python                | Various                    | F        | F        | F        | Y          |
| MITMf       | 2015-08-28            | Python                | Various                    | F        | F        | F        | Y          |
| Ettercap    | 2020-08-01            | C                     | Various                    | F        | F        | F        | Y          |
| bettercap   | 2021-05-22            | Go                    | Various                    | F        | F        | F        | Y          |
| mitmproxy   | 2021-01-21            | Python                | HTTPS, WS, TLS, TCP        | P        | P        | P        | Y          |
| mProxy      | Pre-Releases only     | Go                    | MQTT, MQTT/WS              | F        | F        | F        | Y          |
| IOXY        | Source only           | Go                    | MQTT, MQTT/WS, MQTT/TLS    | F        | F        | F        | Y          |

**Table 5.1.:** Comparison of existing software where *R*, *W* and *D* describe read, write and deletion capabilities, respectively. *F*, *N* and *P* indicate full, no or partial functionality, respectively. *Y* and *N* indicate whether a tool is released under an OSS license or not.

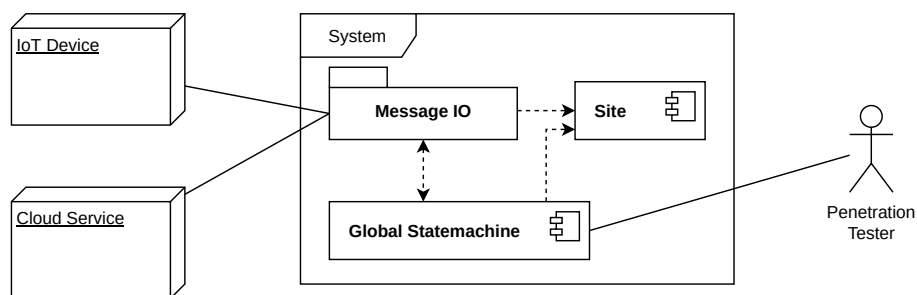
## Chapter 6

# Conceptual Design

Building on the software design of the first prototype presented in section 5.1.3 and the insights gained in section 5.1.6, two design concepts were worked out. The following sections will detail components and principles of both concepts.

### 6.1. Design #1: Monolithic Proxy Application

This design concept is based on the general ideas presented in section 5.1.3 (e.g. state-machines, network stacks and pipes) and employs a basic architecture shown in figure 6.1.




**Figure 6.1.:** High-level component diagram of the proxy application concept

#### 6.1.1. High-level Overview

As discussed in the previous chapter, the requirement “F2 Network Stacks” introduces the need for dynamically initialized objects which in this concept is imple-

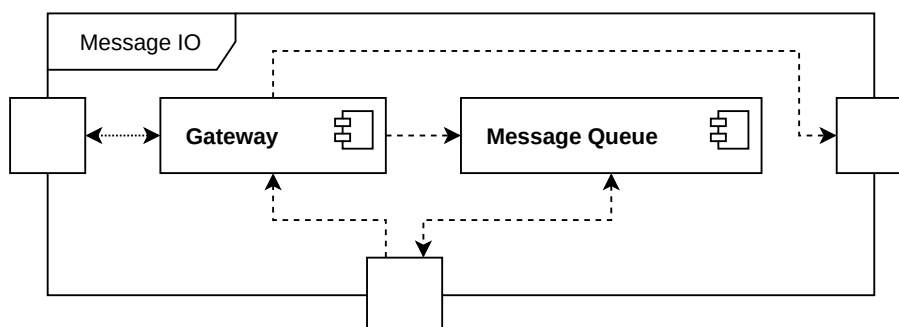
## 6. Conceptual Design

mented by making use of the abstract factory pattern in the “Site” component (depicted in figure 6.2). This component allows for registering *Factories* that are used to initialize objects. Similar to the implementation in the first prototype, factories initialize objects using metadata supplied from a configuration file.



**Figure 6.2.:** A simple variation of the abstract factory pattern. Contrary to the design of the first prototype in section 5.1.3, this variation does not specify the abstract type of the products that are built as return types but as part of the meta data used for object creation.

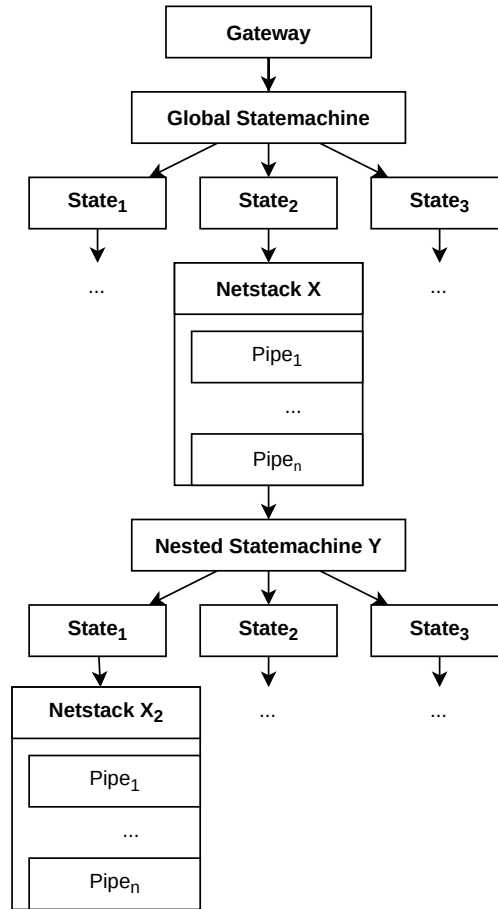
Communication with other systems is encapsulated into the “Message IO”-package shown in figure 6.3. Applications that are tested by penetration testers are connected to sockets provided by the “Gateway” component and temporarily stored in a message queue to be processed by the network stacks organized by the “Global Statemachine”. Similar to the “Server” interface used in the first prototype, gateways provide means of communicating with external systems and receiving and sending messages. They are highly abstract and meant to be used for implementing interfaces for any kind of communication protocols and technologies, such as IP-based TCP and UDP communication but also other protocols such as USB, Bluetooth, ZigBee or KNX. It should be noted that the static view of the design is rather



**Figure 6.3.:** The “Message IO”-package. The white boxes indicate ports for communication with outside components.

simple due to its dynamic runtime behaviour: many instances and relationships are only instantiated at runtime and not pre-determined. A schematic representation of

the dynamic structure and interweaving of state-machines, network stacks and pipes (in this concept called a “*pipeline*”) is shown in figure 6.4. This figure highlights



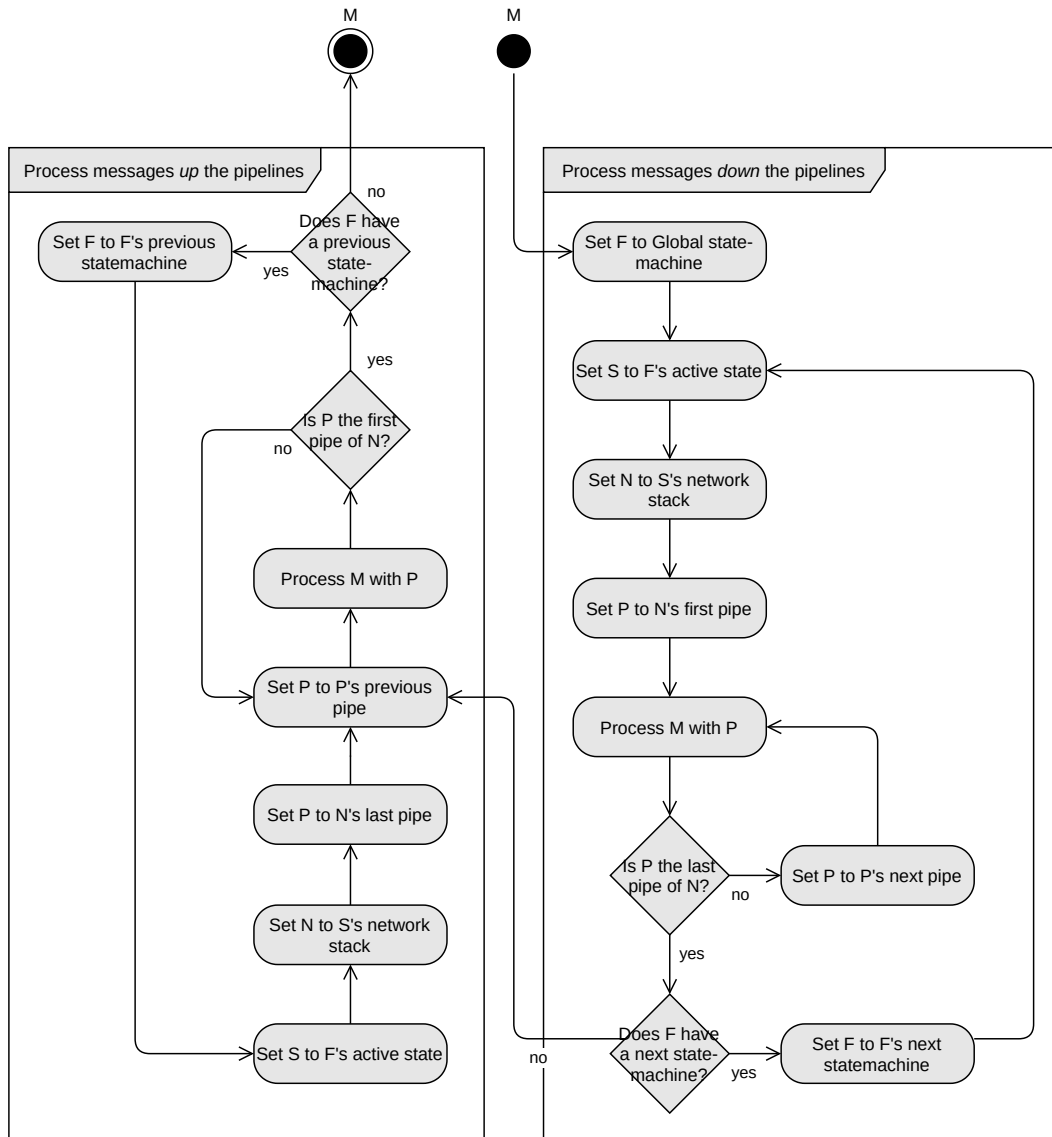
**Figure 6.4.:** An abstract representation of the runtime hierarchy of nested FSMs and network stacks. The active states of the FSMs determine which network stacks are used, resulting in a chain of FSMs and network stacks, referred to as a “pipeline”.

a series of active state-machines and network stacks which together constitute the active pipeline.

Figure 6.5 illustrates the recursive nature of this concept processing (dequeued) messages:

1. A state-machine  $F$  (initialized with the global state-machine instance) relays messages  $M$  through its active state  $S$ ’s network stack instance  $N$ .
2. In  $N$ , all of its pipes  $P$  process  $M$  until the end of  $N$  is reached ( $P$  does not hold a reference to a succeeding pipe instance). If  $F$  holds a reference to a succeeding FSM,  $F$  is set to this reference and the process continues from step 1.

## 6. Conceptual Design



### Diagram Legend

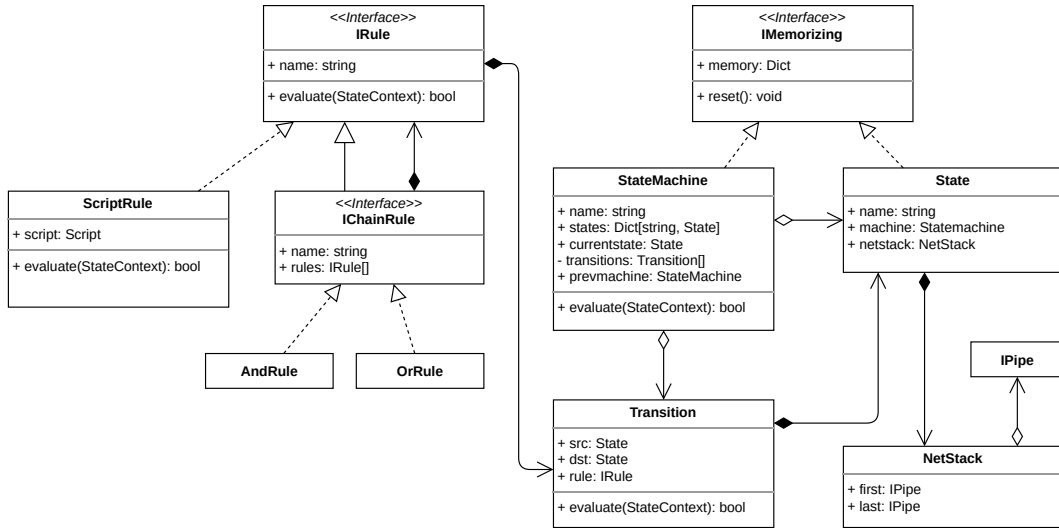
F State Machine  
 S State  
 N Network Stack  
 P Pipe  
 M Message

**Figure 6.5.:** Message processing through an architecture of nested FSMs and network stacks

3. If  $N$  does not hold a reference to a nested FSM, the end of the network stack is reached and the direction of traversing the network stack is reversed.
4.  $P$  is set to  $N$ 's last pipe instance and  $M$  is processed by  $P$  until the start of  $N$  is reached (i.e.  $P$  does not hold a reference to a preceding pipe instance).

If  $F$  holds a reference to a preceding FSM instance,  $F$  is set to this reference,  $N$  is set to  $F$ 's network stack reference and the process continues from step 4.

5. If  $N$  does not hold a reference to a preceding FSM, the beginning of the whole pipeline is reached and  $F$  is the global state-machine.



**Figure 6.6.:** *StateMachines* used rules to determine whether state changes should take place. *IChainRules*, *AndRules* and *OrRules* allowed to combine multiple rules and provided a way to specify logic using *OR* and *AND* operators.

### 6.1.2. State-Machines

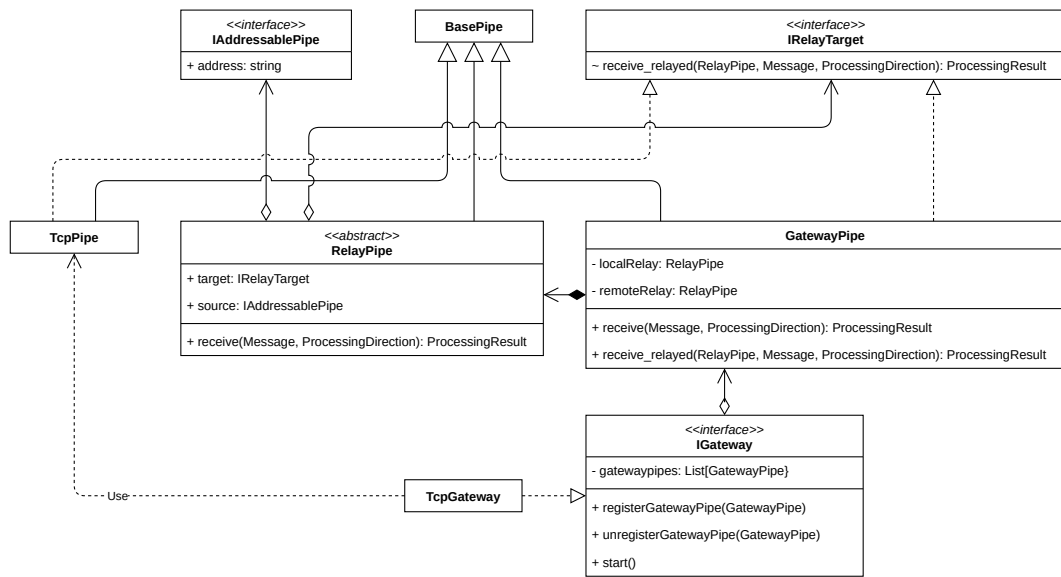
The classes related to the state-machine component are shown in figure 6.6: *StateMachines* hold a set of *States* and *Transitions*. In order to change states, state-machines evaluate a context by checking each of their transitions for whether their conditions for transition are met or not. This context is an aggregation of the *memory* of each state-machine and their active states in the active pipeline. Transitions are defined by a source state, destination state and an *IRule* that evaluates a given context. Rules can be concatenated with logical *AND* or *OR* operators and are designed to be scripts that operate on the given context. This allows the creation of nested rules such as the following one:

$$\text{changeToWS}(c) = \text{AND}(\text{clientUpgrade}(c), \text{serverUpgrade}(c))$$

## 6. Conceptual Design

In this example, a transition with the above rule would evaluate to *true* and trigger a state transition in a state-machine when the aggregated memory *c* of all state-machines and their active states of the active pipeline indicated that an HTTP request was detected that requested an upgrade to the WS protocol (for instance, *clientUpgrade* would look for an entry *clientUpgradeRequested* in *c* and evaluate its contents) and that an HTTP response was detected that confirmed the upgrade request. This would allow a state-machine to detect upgrades of HTTP communication to the WS protocol.

States hold a “NetStack” which in turn encapsulate a series of connected pipes, holding references to this series’ first and last elements.



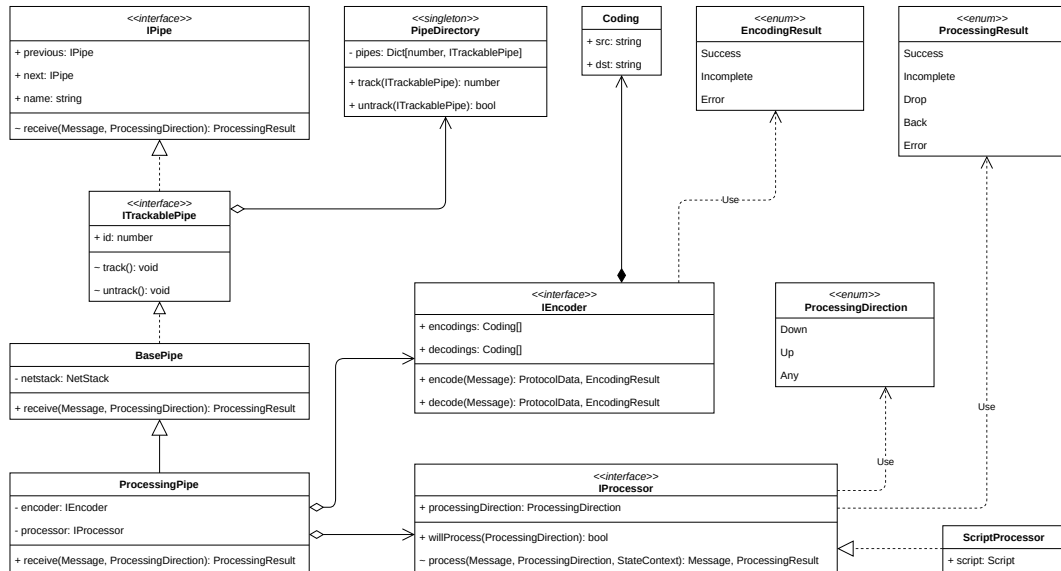
**Figure 6.7.:** By introducing *Gateways*, *GatewayPipes* and *RelayPipes*, multiplexing pipes could be implemented.

### 6.1.3. Gateway

The gateway component is defined by the “IGateway” interface shown in figure 6.7. It is designed to be run as a service in a separate thread that interacts with communication interfaces on machines (i.e. Bluetooth dongles or Ethernet interfaces). During operation it accepts incoming connections  $C_I$  and creates its own respective outgoing connections  $C_O$  to remote servers. Pairs of connections  $C_I$  and  $C_O$  are held in technology-specific pipe implementations and encapsulated in individual “RelayPipe” instances. Those RelayPipe instances are assigned to “Gate-



wayPipe” instances. Improving on the first prototype’s design, the GatewayPipe acts as a multiplexing pipe that accepts messages originating from the two encapsulating “RelayPipes” that act as two communication ports (e.g. the client device and the cloud server of scenario #2 described in section 5.1.1) that hold information about the address of their communication peers in their address field (e.g. an IP address of the remote peer). For instance, two TCP client sockets can be handled by two “TcpPipes” (that inherit from the RelayPipe class), allowing TCP packets to be routed into the pipeline via a GatewayPipe. Once messages are processed and sent back up the pipeline to a GatewayPipe, the GatewayPipe can find the correct RelayPipe to relay the message to by comparing their addresses with the message’s “MessageDirection” information.



**Figure 6.8.:** The interfaces, classes and enums used to represent pipes, their specializations and associated classes. This iteration separates *BasePipes* from *IEncoders* and *IProcessors* so *BasePipes* only implement routing of messages.

#### 6.1.4. Pipes

Building upon the approach of routing and processing messages via pipes discussed in section 5.1.3, this design concept addresses some inconsistencies of the former design and adds needed flexibility. As shown in figure 6.8<sup>1</sup>, the “IPipe” interface is reused from the first prototype and extended by the “ITrackablePipe” interface that adds a unique identifier to pipes. This enables the application to easily locate pipes


<sup>1</sup> A larger printout A.1 can be found in the appendix.

by looking up their identifiers in the “PipeDirectory”, allowing to interact with and inject messages into individual pipes directly.

A “BasePipe” implements the ITrackablePipe interface as well as simple routing logic for forwarding messages up and down pipelines. However, only “ProcessingPipes” actually perform any kind of operations on messages directly: they can employ “IEncoders” for (de-)serialization and “IProcessors” for transformation of messages. Contrary to the design concept of the first prototype, IEncoders need to specify which data formats they support as source and target encodings. This allows the implementation of multiple IEncoders for the same protocol that work with different source or target data formats. For example, some IEncoder may only provide decoding functionality for raw binary data into HTTP messages with raw binary bodies while another implementation provides functionality to encode strings into HTTP message bodies. In the first prototype’s design, the very concept of filters was only vaguely described and lacked a clear and concise interface. This issue is resolved in this next iteration of the design concept:

- Filters are renamed to “IProcessors” (conveying the purpose and meaning of the interface in its name).
- IProcessors specify a “ProcessingDirection” that determines whether messages shall be processed on their way *down* or *up* a pipeline or in any direction, effectively granting control over applying transformations on messages. This can be helpful when transformations shall only be applied in one direction or maybe only once in a pipeline, like replacing the contents of the body of an HTTP message.
- An IProcessor can apply logic to messages in its *process* method that also receives the pipeline’s context. The returned “ProcessingResult” indicates success or failure of the operation or whether the IProcessor requests dropping a message or sending it back up the pipeline.

While there are many opportunities for specific implementations of the IProcessor interface, one general implementation is envisioned by the design concept: a simple “ScriptProcessor” allows penetration testers to supply scripts that are executed at runtime and allow transformation of messages. This directly fulfils the requirement “F5 Scripting”.



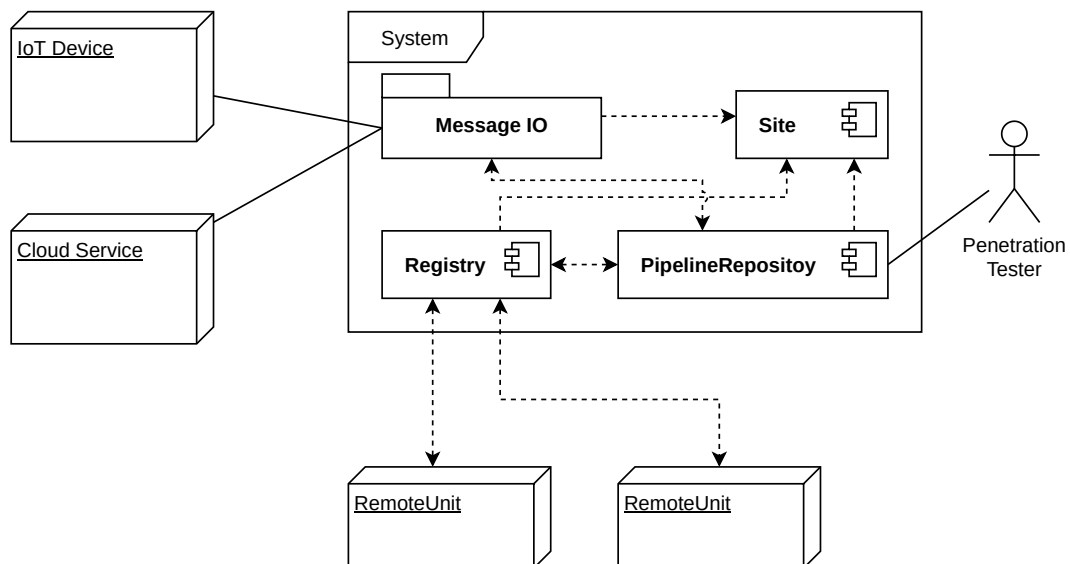
**Figure 6.9.:** The data-structures used to represent messages, their meta-data and payloads.

### 6.1.5. Messages

Compared to the first prototype's design, the data-structures that represent messages are mostly unchanged (as shown in figure 6.9). However, during implementation and testing of the first prototype it became apparent that in some cases “historical” information about messages was required. This iteration of the design adds a list of “ProtocolData” instances to messages that specify information about the protocol, message headers and the payload (“ISource”). IProcessors and IEncoders append newly transformed or (de-)serialized ProtocolData instances to messages. So over time, a message contains records of all those operations performed on it. This information can be useful in a number of cases like serialization: when a message is deserialized (e.g. the payload of a WS message is extracted) on its way down a pipeline, important information about the formerly encapsulating protocol is lost (such as the WS frame's flags). When a message is serialized on its way back up a pipeline, an IEncoder would have to generate this information or try and deduce it from the message, which is not always possible. However, since it can access the messages' history and former ProtocolData, it can read the original information and use this for serialization.

## 6.2. Design #2: Distributed Proxy Services

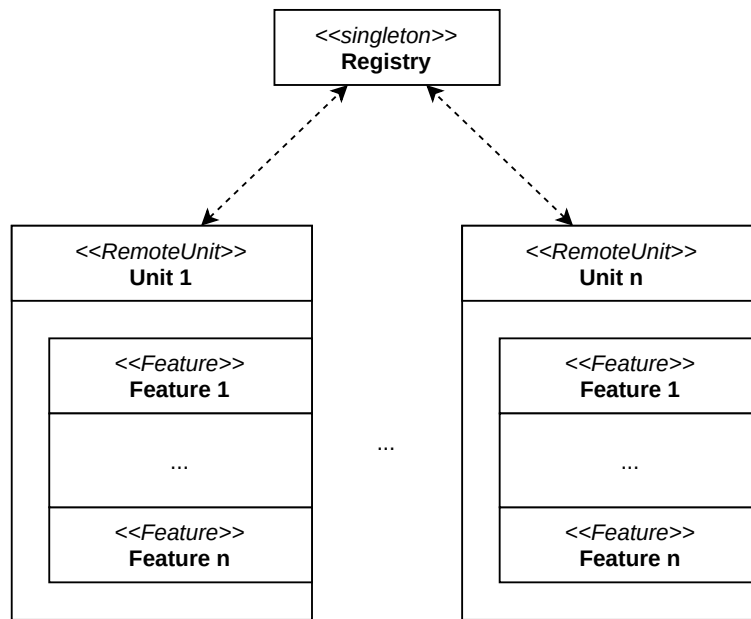
The design shown in section 6.1 was an iteration of the design worked out for the first prototype in section 5.1.3 and addressed some fundamental, architectural flaws and aimed for better flexibility and more meaningful interface definitions. However, it did not address other problems that were encountered during the implementation of the first prototype: constraints in platform, framework and programming language compatibility and flexibility. As a consequence of these constraints, the proxy application needed to be developed as a monolithic application. Therefore, each extension, like additional IEncoders that added support for new protocols, was required to be implemented in the same programming language and run on the same platform and as part of the same process as the proxy application. This also effectively limited the available selection of libraries. Another potential problem of the former design concept was the tight coupling of pipes and the deeply nested structure and hierarchy of state-machines, pipes and network stacks. While this architecture allowed to implement routing messages through composition (*by design*), it greatly added to the runtime complexity and made debugging the application significantly harder.



**Figure 6.10.:** A component diagram showing the external systems communicating with the system and their connection to internal components.

### 6.2.1. Overview

Another iteration of the design (shown in figure 6.10) was made to address these issues. While the “Message IO” and “Site” components are left unchanged, the global state-machine is replaced by the “Registry” and “PipelineRepository” components that allow for de-centralized and more controlled processing of messages.

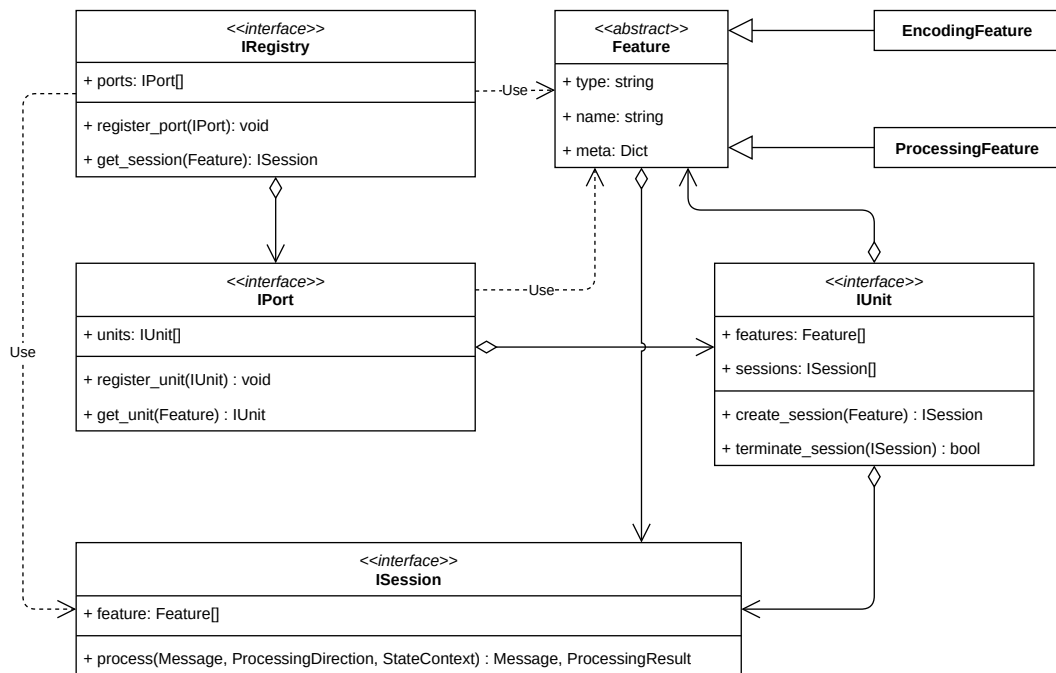


**Figure 6.11.:** In this concept,  $n$  separate *Units* can be registered at a central *Registry*. *Units* may provide an arbitrary amount of *Features*.

### 6.2.2. Registry and Units

To implement de-centralized processing of messages, a central registry is required to register remote units (shown in figure 6.11). As can be seen in figure 6.12, the central registry is represented by the “IRegistry” interface that allows remote units to register themselves and allows the PipelineRepository to request sessions to units that implement requested features. Remote units can be remote machines that implement the “IPort” interface that provides a list of “IUnit” instances. An IUnit implements one or more “Features”, such as specific (de-)serialization or other processing, and effectively provides IEncoder and IProcessor functionalities. This transforms formerly direct calls to IEncoders and IProcessors to Remote Procedure Calls (RPCs). Since IEncoders and IProcessors can be stateful, IUnits initialize them in “ISessions” for each requested feature. The IRegistry and IPort interfaces

## 6. Conceptual Design



**Figure 6.12.:** Visualized in this diagram, the distributed implementation and central registration of features is the core idea of the Registry and Unit components.

are explicitly kept rather simple and unspecific to the exact means of communication between them so that they can be implemented in various ways, making use of various Inter-Process Communication (IPC) techniques.

Also, to transmit data between the proxy application and its remote units, this data needs to be (de-)serialized and a format for serialization has to be chosen.

### 6.2.3. PipelineRepository

The PipelineRepository component holds information about all configured pipelines (that is a flattened representation of the hierarchically configured state-machines and network stacks) and their contexts. This allows to remove the pipes from the software architecture, providing better traceability of messages throughout the system and makes debugging the high-level application logic more accessible. Also, organizing network stacks and state-machines in one central place encourages creation of means to interface with these mechanisms such as REST-APIs that let penetration testers inspect the message queue and ongoing processes.

#### 6.2.4. State of the Design Concept

This design concept promises to solve severe issues of the previous design iteration shown in section 6.1 and already defines some very high-level components. However, due to time constraints some components' designs were not finished and require further work on specifics. For this, certain questions need to be answered and translated into the design:

- **PipelineRepository:** How exactly is the hierarchy of state-machines and network stacks flattened? How is this flattened hierarchy represented in data-structures? How are instances of individual state-machines and network stacks initialized and organized for individual gateway-connections?
- **ISession:** How is information relevant to IEncoder and IProcessor instances (such as ScriptProcessors' Script instances) passed to remote units?
- **IRegistry/IPort:** How exactly are RPCs performed? Are there mature and appropriate frameworks that provide RPCs implementations?

#### 6.2.5. Comparison of Both Designs

Both design concepts discussed in the previous sections, the monolithic and the distributed concept, promise to solve specific problems. The following paragraphs compare both concepts on the base of a set of core design aspects:

**Software architecture** The monolithic concept suggests a centralized and self-contained design that combines the high-level business logic of a proxy-application with low-level tasks such as (de-)serialization of various protocols. It is designed to be run on a single machine.

In contrast to this, the distributed concept separates the high-level business logic (like routing messages) and low-level tasks as part of a client-server model: the high-level logic is implemented in the central proxy server while low-level tasks are isolated into separate remote services. These services can either be run on the same machine the proxy server runs on or on external machines. Through dynamic creation and registration of remote service instances, this concept also implements scalability. Since there is no restrictions to the programming languages, platforms

or frameworks used by remote services, the concept also embraces platform compatibility.

**Complexity** The reliance on deeply nested data-structures has a high impact on the complexity of the monolithic concept at runtime. This makes debugging an implementation of this concept significantly harder and more time-consuming. However, adding new extensions to this concept becomes a comparatively easier task as integration of such new extensions only takes place on a source-code level.

As opposed to this, the distributed approach simplifies the high-level tasks such as routing messages by introducing distributed components that allow traceability and thus establish transparency. The offloading of protocol implementations into logical units that are accessed via IPC however contribute to a more challenging deployment of the application. Also, due to their distributed nature, debugging these remote units can introduce further problems (e.g. connection losses and high latencies) and requires a more sophisticated and complex testing environment than the monolithic concept.

**Maturity** Some of the core components of the monolithic concept were already tested and proven by the first prototype. For instance, linked pipes proved to be an effective means to route and process messages up and down a processing pipeline. However, due to time-constraints, other core concepts such as the state-machines could not be tested. Regarding completeness, it is noteworthy that this concept's interfaces are well-defined.

Contrary to this, the distributed concept is not finished and requires further work to clear up a number of essential questions before it can be completed. Also, the previously proven effective idea of using pipes for routing and processing is removed in the distributed approach. This denies the approach any effectiveness acquired through previous design iterations.



## Chapter 7

# Implementing the Modular Proxy Application

This chapter covers an exemplaric implementation of the monolithic concept that was worked out in section 6.1, starting with formally describing the goals and constraints of this implementation in section 7.1. Afterwards, an overview and comparison of available and suitable tools for the task is performed in section 7.2. The chapter concludes with details about the implementation of individual components in section 7.3, describing how specific challenges were overcome.

### 7.1. Goals and Constraints

The goal of this thesis' implementation was to implement the “Monolithic Proxy Application” design concept described in section 6.1 to a maturity level that allowed to test its usefulness and effectiveness in the testbed described in section 5.1.4 that aimed to represent scenario #2 “IoT Cloud Application” discussed in section 5.1.1. Thus, a focus was set on implementing a vertical prototype that featured important core components (such as factories, state-machines and network stacks) and a set of exemplaric protocol implementations (HTTP, WS and MQTT). Similar to the first prototype discussed in section 5.1, this prototype was a proof-of-concept implementation and neglected quality attributes such as usability and performance. The prototype had the working title “net-riot”, which indicated that this was a net-working tool and was to be used in the IoT context.

### 7.2. Tool Selection

To choose the tools for implementing the design concept, a list of requirements for tools was inferred from the software requirements discussed and expert interviews shown in chapter 5:

- T1 Scripting:** The tool must provide scripting capabilities that allow penetration testers to execute complex scripted operations on messages.
- T2 Libraries:** In order to avoid custom implementation of the HTTP, WS and MQTT protocols, the tool must provide a rich set of libraries that can be used to work with said protocols.
- T3 Deployment:** To allow the prototype to be installed in an uncomplicated way, the tool must provide or support mechanisms that simplify deployment, such as code compilation and static linking of dependencies or containerization.
- T4 Accessibility:** The tool must be powerful and complex enough to solve the software requirements and implement the design concept, but it must also feature a “barrier or entry” that is low enough so extending the application is feasible for open source developers.

Regarding the selection of a suitable programming language, it was found that Python satisfied all of these requirements:

- Python is a free, open source and general purpose programming language that was first released in 1991 and is being continuously improved and updated.
- The built-in *exec*<sup>1</sup> function provides execution of arbitrary Python code at run-time. Although it is infamous for its security implications, it is very suitable for scripting.
- The Python Package Index (PyPI) is a public repository of more than 300.000<sup>2</sup> Python packages that can be installed using the *pip* command-line tool. There are numerous packages that implement the protocols WS, MQTT and HTTP.
- Python supports multiple ways to deploy projects, including packaging projects into executable files<sup>3</sup> and containerization<sup>4</sup>.

---

<sup>1</sup><https://docs.python.org/3/library/functions.html#exec>

<sup>2</sup>Based on PyPI's statistics: <https://pypi.org/>

<sup>3</sup><https://packaging.python.org/overview/#bringing-your-own-python-executable>

<sup>4</sup>e.g. using Docker [https://hub.docker.com/\\_/python/](https://hub.docker.com/_/python/)

- Its comparatively simple syntax and its design philosophy that values accessible code higher than performance<sup>5</sup> encourage readability and maintainability in Python projects. When implemented, this makes Python an accessible programming language. Also, its optional static typing allows to omit redundant type information for simple methods and to provide explicit type information for complex and shared pieces of code like algorithms and interfaces.

Git was used for version-control and Microsoft Visual Studio Code was the Integrated Development Environment (IDE) used for implementation.

### 7.3. Individual Components

The following sections discuss especially challenging aspects of net-riot's implementation, which problems were encountered and how they were solved.

#### 7.3.1. Gateways

The protocols used in scenario #2 (HTTP, WS and MQTT) are used on top of TCP. Therefore, net-riot implemented a TCP-Gateway that allowed it to create TCP server sockets to listen on for incoming connection requests. For incoming connections  $C_I$ , respective outgoing connections  $C_O$  were initialized and connected to a preconfigured remote server. For each of those connections, "TcpPipe" instances  $P_I$  and  $P_O$  were initialized that ran in separate threads and accepted incoming packets. The gateway then initialized two "TcpGatewayPipe" instances with  $P_I$  and  $P_O$  that routed messages originating from the TcpPipes into the pipeline and from the pipeline to the correct TcpPipe instance.

Messages that originated from gateways were temporarily stored in a queue so that only a single message was processed in the pipeline at any given time. It was found that if multiple messages were processed simultaneously, the global state-machine could change states while a message was still being processed in a then inactive state. This resulted in this state's network stack being left unconnected and unable to route the message back up the pipeline. For the same reason, net-riot only supported one single connection.

---

<sup>5</sup>Described in 19 aphorisms in "The Zen of Python" (<https://www.python.org/dev/peps/pep-0020/>)

### 7.3.2. Encoders

For each of the protocols used in scenario #2, net-riot implemented a separate IEncoder.

**HttpEncoder** For HTTP, no library was found that parsed HTTP requests or responses into low-level representations that did not discard essential information. However, since HTTP is a comparatively simple, text-based and stateless protocol, a custom IEncoder was implemented that parsed HTTP requests and responses from raw binary data and allowed to assemble requests and responses from processed messages. For assembly, the implementation used the HTTP header information contained in a message's "history" field that held the headers that were parsed when the message was first parsed by this encoder. One practical pitfall of the custom implementation was the "Content-Length" HTTP header that indicated the number of bytes contained in an HTTP request body or HTTP response body. Systems that parse HTTP messages (such as web-servers and browsers) use the value of the "Content-Length" header to read the indicated number of bytes from a TCP stream and associate it with the parsed headers. If the length of a message body is modified by a proxy, the "Content-Length" header indicates the wrong number of bytes to read: this will either result in reading too few bytes (thus, discarding information) or reading too many bytes rendering future requests malformed. To solve this issue, the custom encoder provided the configurable flag *recalculateContentLength* that, if set to true, dynamically calculated the value of the "Content-Length" header to reflect the actual length of bytes contained in the message's BinarySource.

**WsEncoder** The library used for WS implementation was "websockets"<sup>6</sup>. It offered methods for parsing (*framing.Frame.read*) and assembling (*framing.Frame* constructor) WS frames. These functions worked fine for regular WS frames. However, to save bandwidth, some WS client and server implementations make use of PMCE. The use of PMCE is indicated by the "rsv1" bit of WS headers being set to true. While the library implemented PMCE (*extensions.permmessage\_deflate.PerMessageDeflate*) and using the extension worked on the first messages of a WS connection, it would fail to correctly compress messages after it processed a number of messages, causing the remote WS clients and servers to terminate the connection. It was found that the PMCE implementation was stateful and not reset after being

---

<sup>6</sup><https://github.com/aaugustin/websockets/>, commit 6b5cbaf41cdbc9a2074e357ccc613ef25517dd32

used. This issue was solved by initializing new *PerMessageDeflate* instances per WS frame assembly.

**MqttEncoder** For (de-)serialization of MQTT messages, the library “hbmqtt”<sup>7</sup> was used in net-riot. The library implemented individual classes for each MQTT message type and provided methods for parsing (*from\_stream*) and assembly (*to\_bytes*) of packets. However, the library did not provide any method that chose the correct class for (de-)serialization of a given binary buffer or processed message. Therefore, this functionality was implemented in net-riot: the *MqttEncoder* defined a dictionary that mapped MQTT message types to a tuple of classes (*Packet*, *VariableHeader* and *Payload*) provided by the library and attempted to parse the basic MQTT message header of a binary buffer. This header included the MQTT message type that was then used to look up the correct classes for parsing in the dictionary. When MQTT messages were re-assembled by the *MqttEncoder*, it extracted the message type from the parsed header (that was stored in the message’s “history” field) and used this to look up the correct classes for serialization in the dictionary.

### 7.3.3. Scripting

In order to allow execution of scripts on messages, net-riot implemented the “Script-Processor” discussed in section 6.1. This *IProcessor* implementation used Python’s built-in *exec*<sup>8</sup> function for executing arbitrary Python code (shown in listing 7.1). The function allows the caller to define the available objects of the local and global scopes in the called code. This mechanism was used to pass the message instance, processing direction, context and “result” object to the script in the local scope. The *ProcessingResult* values were passed to the script in the global scope which allowed to script to assign one of these values to the “result” object in the local scope. After execution, the “result” value was evaluated and returned by the *ScriptProcessor*. This effectively enabled scripts to:

- Read and write fields and payloads of messages.

---

<sup>7</sup><https://github.com/beerfactory/hbmqtt/>, commit 31165fb0e827925417f99a7b1f475a9d67e1c72f

<sup>8</sup>Python’s *exec* function and similar functions of other programming languages and frameworks that allow code execution are infamous for code injections and pose a security risk. However, this risk was tolerated for this prototype as it is designed and intended to be run in a controlled environment.

## 7. Implementing the Modular Proxy Application

---

- Read and write the active pipeline’s context and thus, its state-machines’ memory. This could be used to trigger state transitions.
- Control whether messages were processed further, being immediately sent back up the pipeline, being dropped or ignored.

```
class ScriptProcessor(IProcessor):
    def __init__(self, script: Script, inDirection: ProcessingDirection =
        ProcessingDirection.ANY):
        super().__init__(inDirection)
        self._script_ = script

    def process(self, msg: Message, direction: ProcessingDirection, context:
        StateContext) -> Tuple[Message, ProcessingResult]:
        _locals = {'msg': msg, 'direction': direction,
            'res': ProcessingResult.SUCCESS,
            'context': context}
        _globals = {
            'SUCCESS': ProcessingResult.SUCCESS,
            'ERROR': ProcessingResult.ERROR,
            'INCOMPLETE': ProcessingResult.INCOMPLETE,
            'DROP': ProcessingResult.DROP,
            'BACK': ProcessingResult.BACK,
        }
        try:
            exec(self._script_.body, _globals, _locals)
        except Exception as e:
            print(f'Script error: {str(e)}')
            return (_locals['msg'], ProcessingResult.ERROR)
        return (_locals['msg'], _locals['res'])
```

**Listing 7.1:** The ScriptProcessor implementation of net-riot using Python’s built-in exec function.

For net-riot to support MQTT communication that was tunnelled via WS it had to detect upgrades of HTTP connections to WS. This was achieved by ScriptProcessor instances that executed scripts that examined the HTTP communication and set values in the memory of specific state-machines using the context that was supplied to them. Listing 7.2 shows the script that detected a server’s HTTP response to an upgrade to WS: if there was an “Upgrade” header with the value “websockets” and if the status code of the response was 101, the key “serverUpgrade” of the nested state-machine “http” was set to *True*. If any of these checks failed, it was set to *False*. This script worked in conjunction with another script that performed similar checks to detect a client’s request to upgrade the connection. The “http” state-machine used “ScriptRules” that evaluated the value of the “serverUpgrade” key and triggered a state transition if both a client’s upgrade request and the server’s upgrade response were detected.

```
def process(msg, context):
```

```

_upgrade = msg.latest_data.fields.get('Upgrade')
_status = msg.latest_data.fields.get('::Status')
_upgradeToWs = _upgrade.lower() == 'websocket' if _upgrade else False
_status101 = _status.find('101') != -1 if _status else False
context.memory['http_to_ws']['http']['serverUpgrade'] = _upgradeToWs and
    _status101
return SUCCESS

res = process(msg, context)

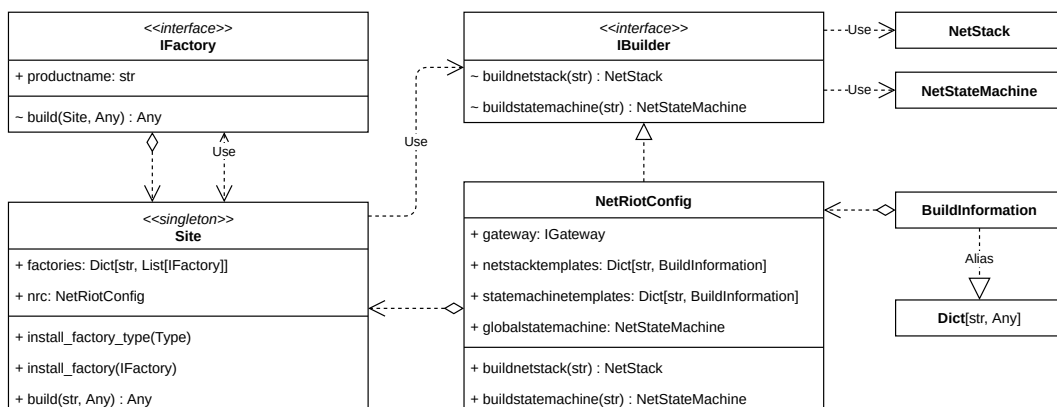
```

**Listing 7.2:** The script net-riot used to detect upgrades of HTTP connections to WS.

### 7.3.4. Configuration Parsing and Building

The requirement “F2 Network Stacks” mandated the capability to read a configuration file at runtime and initialize the therein specified hierarchy of state-machines and network stacks.

**JSON Configuration and Schemas** For formal specification and representation of state-machines and network configurations, net-riot made use of the JSON format. In over 600 lines of code, a JSON schema was defined that specified the structure and types of net-riot configuration files. JSON and JSON schema were chosen over alternatives such as XML and YAML due to their simple syntax, flexibility and powerful features (such as *definitions*).



**Figure 7.1.:** The implementation of the variant of the abstract factory pattern discussed in section 6.1. The Site component and IFactory instances recursively call each other to create and composite object instances.

**Factories, Builders and Templates** The “Site” component was briefly discussed in section 6.1 and envisioned as an implementation of the abstract factory pattern. Figure 7.1 shows net-riot’s implementation of the Site component and its associated interfaces and classes. The Site component holds a dictionary that maps product-names to factories. In turn, factories can instantiate objects such as pipes, gateways and state-machines (each differentiated by distinct names). Listing 7.3 shows a simple factory implementation that is used to instantiate “ScriptRule” objects. However, prior to calling the constructor of the ScriptRule class, it needs to acquire a “Script” instance. It does so by requesting the Site to build a Script instance with the supplied information. In turn, the Site looks up the correct factory to use for instantiating the requested Script instance and calls it. Thus, factories can recursively request further objects to be built by the Site while they themselves provide the functionality to produce a single product each.

```
class ScriptRuleFactory(IFactory):
    def __init__(self) -> None:
        super().__init__(RULE_SCRIPT)

    def build(self, plant: Site, obj: Dict[str, Any]) -> Any:
        _scriptobj = obj['config']
        _name = obj['name']
        _script = plant.build(SCRIPT, _scriptobj)
        return ScriptRule(_name, _script)
```

**Listing 7.3:** A simple factory that requested the build of a “Script” instance to instantiate a “ScriptRule”.

The “NetRiotConfig” class implements further logic to net-riot that allows to dynamically instantiate whole state-machines and network stacks. It does so by saving the configured state-machines and network stacks as *templates* and passing them to the Site to build. For instance, net-riot’s configuration file may specify a state-machine by the name of “http\_to\_ws” that is specified to be used as a nested FSM of a state “entry” of the global state-machine. When the global state-machine is built and an instance of the nested FSM is requested, the Site component calls NetRiotConfig’s *buildstatemachine* method and supplies the name “http\_to\_ws”. In turn, NetRiotConfig calls the Site component’s *build* function and supplies the template of the requested state-machine.

Also, net-riot used Python’s “inspect”<sup>9</sup> package to dynamically acquire all types of implemented factories and register them in the Site instance (shown in listing 7.4).

```
import inspect
import parsing
```

---

<sup>9</sup>Python’s inspect package provides functionalities that allow examination of code at runtime, similar to “reflection” packages present in other programming languages such as Java and C#.



```
from py_linq import Enumerable

factory_types = Enumerable([m for _, m in inspect.getmembers(parsing)]\
    .where(lambda m: inspect.isclass(m) and not inspect.isabstract(m) and parsing.\
        factory.IFactory in inspect.getmro(m))\
    .to_list())

for factory_type in factory_types:
    parsing.SITE.install_factory_type(factory_type)
```

**Listing 7.4:** Dynamic registering of all factory implementations.



## Chapter 8

# Evaluation and Lessons Learned

This chapter attempts to evaluate the work the proxy application and discuss what challenges were encountered and what were the lessons learned from those. The project timeline will allow a quantitative overview of the project progression and show what parts of the project slowed down progress. Then, an overview of the qualitative aspects of the deliverables will discuss the maturity of the design concepts presented in chapter 6 and the implementation of the monolithic design in chapter 7 and which parts reached a satisfactory level. Lastly, this chapter will present the lessons learned during this work.

### 8.1. Quantitative Overview: Time Management

Comparing the planned thesis schedule to the actual course it has taken, this section discusses how the intended plan was implemented and changed at certain places. Also, it will examine the causes of the delays during development.

#### 8.1.1. Project Timeline

Table 8.1 shows the initially planned thesis schedule divided into four phases, laying out the course of the thesis over a span of 24 weeks.

**1. Preparation** The initial phase covered preparation tasks for further work on the thesis. Literature research on the topics covered and touched in this thesis was

## 8. Evaluation and Lessons Learned

---

| Phase / Task             | Duration          |
|--------------------------|-------------------|
| 1. Preparation           | 4 weeks (16, 66%) |
| Literature Research      | 1 week            |
| Expert Interviews        | 1 week            |
| Testbed Configuration    | 2 weeks           |
| 2. Prototype             | 7 weeks (29, 16%) |
| Prototype Conception     | 2 weeks           |
| Prototype Implementation | 4 weeks           |
| Expert Feedback          | 1 week            |
| 3. Functional Prototype  | 7 weeks (29, 16%) |
| FP Conception            | 2 weeks           |
| FP Implementation        | 4 weeks           |
| Expert Feedback          | 1 week            |
| 4. Finalization          | 6 weeks (25%)     |
| MQTT Case Study          | 2 weeks           |
| Thesis Finalization      | 4 weeks           |
| <i>Total</i>             | <i>24 weeks</i>   |

**Table 8.1.:** Initially planned schedule for the thesis

carried out. Related work on IoT and ICS security analysis (as discussed in chapter 2) was of special interest as those showed what approaches had been taken to assess security implementations. Also, a testbed (discussed in section 5.1.4) for running the proxy application was built. A decision was made against conducting expert interviews before implementing a first prototype on the assumption that practical experience with the subject matter would benefit the expert interviews. The fact that a number of important questions arose from work on the first prototype later proved this decision to be correct. Performing the literature research and building a testbed was completed within the intended schedule of three weeks.

**2. Prototype** In the second phase, the prototype discussed in section 5.1 was designed and implemented in weekly sprints. Preceding these sprints, a rough design of the prototype's architecture and runtime behaviour was worked out in one week that would serve as a base for further design refinement and implementation in the sprints. These sprints ran for eight weeks in total: the initial design turned out to be too oversimplified so that sprints aiming to design and implement specific components were conducted rather isolated from other components that still needed to be worked on. As a result, both the integration of individual components and their interaction would fail and require redesigns and time-consuming adjustments to their implementation. Also, neither was the prototype mature enough to be used as a proxy application, nor was the resulting design and implementation clean enough to suggest putting further effort into working on them. After these eight sprints, work on this prototype was stopped and the expert interviews discussed in section 5.2 were prepared and conducted. It was found that the project was technically challenging and more complex than initially anticipated so the expert interviews were conducted to aid in re-engineering the design concept.

**3. Functional Prototype** The third phase was intended to yield a design concept mature enough to both fulfil realistic requirements to a proxy application and be implemented. This was initiated by switching the technology stack from TypeScript to Python and re-designing and re-implementing large parts of the first prototype. In order to avoid the same mistake of refining a vague design concept and spending time adjusting the design and implementation to make them work, two weeks were spent on iterations of new design concepts discussed in chapter 6. These concepts did not only define single components but also interfaces that specified how those

## 8. Evaluation and Lessons Learned

| Phase / Task             | Duration        |
|--------------------------|-----------------|
| 1. Preparation           | 3 weeks (12%)   |
| Literature Research      | 1 week          |
| Testbed Configuration    | 2 weeks         |
| 2. TypeScript Prototype  | 10 weeks (40%)  |
| Prototype Conception     | 1 week          |
| Prototype Implementation | 8 weeks         |
| Expert Interviews        | 1 week          |
| 3. Python Prototype      | 12 weeks (48%)  |
| RC Conception            | 2 weeks         |
| RC Implementation        | 10 weeks        |
| <i>Total</i>             | <i>25 weeks</i> |

**Table 8.2.:** Actual schedule of the project

components interacted with each other, aiming for clear separation of components and high flexibility in implementation. Components of the prototype that were independent of the communication protocols used at runtime, such as NetStacks and FSMs, were implemented first over the span of four weeks. Then, implementations for supporting the HTTP, WS and MQTT protocols followed over a span of another six weeks. Work on this prototype was stopped after those ten weeks as the technical difficulties discussed in section 8.1.2 made estimations over the remaining time needed to finish the prototype both hard to make and rather unreliable.

**4. Finalization** The final phase was intended to conduct a case study on how the proxy application would perform on scenario # 2 from section 5.1.1. Tests were made to run the proxy application in the testbed shown in section 5.1.4 which featured the same communication protocols that were used in scenario # 2. However, the proxy application failed to reliably transmit or encode the messages sent between the MQTT client and broker, thus resulting in a broken communication channel. The complex runtime behaviour and very time-consuming debugging of the proxy application (further elaborated on in section 8.1.2) lead to the decision to stop the project.

Table 8.2 shows the actual schedule of the thesis. As can be seen, 88% (22 weeks) of the time working on the thesis was spent designing and implementing the prototypes compared to a planned portion of roughly 60% (14 weeks).

### 8.1.2. Development Challenges

There was a series of development challenges that slowed down implementation of both prototypes considerably:

**Complex runtime behaviour** The combination of nested FSMs and pipelines lead to several problems during development. Even comparatively simple scenarios to use the proxy application in required a complete configuration file. This file was made of a global state machine and at least one network stack. This lead to a dynamic and long chain of references at runtime that made tracing back calls and attributing them to specific instances difficult.

Some problems such as a timing problem in the implementation of FSMs were very time consuming to debug: a FSM would change its state when any of its rules was evaluated successfully and indicated a state change. By design, all FSMs of an active netstack would evaluate their rules when a message entered or left any netstack. When a higher-level FSM (e.g. the global state-machine) changed its state *while* a message was still being processed in a lower-level FSM, the higher-level state-machine would change to another netstack, thus disconnect the lower-level state-machine. Eventually, the message would be processed back up and run into a pipe that had no upstream connection anymore, raising an exception and terminating the program. This particular error was discovered during the implementation and testing of the MQTT encoder, in a runtime setup that involved a global default state-machine with a default TCP netstack and a state-machine that handled HTTP to WS upgrades and processed MQTT messages utilizing network stacks for HTTP and WS/MQTT.

Other problems uncovered design flaws and required prompt changes to the software design or, in some cases, introduced new constraints to the project. One such example was discovered while testing the HTTP encoder implementation using Mozilla FireFox as an HTTP client. When browsing websites, the browser would open multiple connections to the target host to acquire multiple files at the same

time<sup>1</sup>. This required the proxy application to instantiate a new pipeline per incoming connection rather than reside on using a single pipeline. However, the pipeline design dictated that pipes were connected to at most one preceding and one succeeding pipe. In order to connect pipelines to FSMs, the FSMs needed to provide a pipe-interface themselves. Thus, when instantiating multiple pipelines, these could not be connected to the global state-machine because the global state-machine's pipe-interface could only be connected to a single pipeline at a time. Therefore, a multiplexing pipe needed to be implemented to solve this issue. Alternatively, the proxy application could enforce instantiation of one single pipeline only to avoid changes to the software design. For a lab environment, enforcing the use of a single connection might work, however in real scenarios this constraint could potentially lead to the proxy application breaking applications at runtime. To enable future implementations to support multiple connections, the software design was changed in a way that would allow the proxy to handle multiple connections. However, due to its academic nature, the prototype would only support a single connection at a time to reduce complexity.

**Open source libraries** Both prototypes made use of open source libraries that implemented various protocols and included serialization and de-serialization routines for handling protocol specific packets. However, such libraries appeared to be intended to be used for developing applications that used those protocols as a means for transporting data rather than directly parsing packets.

Usually, these libraries would offer an API that allowed to instantiate and operate clients and servers and bind callbacks to events. The implementations of packet serialization and de-serialization were often times hidden through encapsulation, missing typings or poorly documented. For instance, the JavaScript library “ws”<sup>2</sup> provided methods for serialization and de-serialization but lacked typings. Typings for this library were made available by the project “DefinitelyTyped”, however those did not include the classes relevant for serialization and de-serialization (“Sender” and “Receiver”)<sup>3</sup>. At the time of implementing the Python prototype, it used the library “websockets” that offered only an async de-serialization method (“framing.Frame.read”), requiring the use of asyncio which was circumvented by implementing a wrapper around it.

---

<sup>1</sup>For testing single HTTP connections, the key *network.http.max-connections-per-server* could be set to 1 in the *about:config* page.

<sup>2</sup><https://github.com/websockets/ws>, version 7.0.0, commit 092a822a

<sup>3</sup><https://github.com/DefinitelyTyped/DefinitelyTyped/>, commit 4bf23527



The Python prototype also used the “hbmqtt” library to (de-)serialize MQTT messages. The library used an object-oriented implementation for (de-)serializing MQTT messages where a class for each MQTT message type (e.g. *CONNECT*, *CONNACK*...) inherited from an abstract “MQTTPacket” superclass that defined a “to\_bytes” method for serialization and an async “from\_stream” method for (de-)serialization. Since the library did not implement a generic method that parsed a byte-buffer and returned the appropriate MQTT message object, this logic had to be implemented as part of the work in the prototype, requiring investigation of the (largely uncommented) source code of the library as its documentation did not cover these internal (de-)serialization methods but focused on high-level use of the API it implemented.

From a software engineering point of view, omitting public interfaces to internal (de-)serialization methods and forcing specific programming patterns (such as async programming) are perfectly valid decisions in the context of single, individual modules. However, for those reasons, making use of the functionalities implemented in those libraries, was not trivial. It involved developing workarounds and investigating the libraries’ source code which in turn took up time during the implementation phases.

Then there were also instances of incomplete documentation: the Python library “websockets” implemented (de-)serialization of WS packets and also implemented the PMCE of the WS protocol. Calling the (de-)serialization methods of the “websockets” library and specifying the use of PMCE, the first incoming and outgoing messages would be compressed correctly, however following messages would be compressed incorrectly. This rendered the prototype useless as WS may use PMCE by default to reduce bandwidth. The library failed to raise exceptions or return error codes so from the prototype’s runtime point of view it appeared to work just fine. After investigating the library’s source code it was found that the instances implementing the extension were stateful. When supplying newly created instances of said extension implementation to the (de-)serialization methods, they worked as intended, compressing and decompressing any amount of WS packets. This could be due to a multitude of reasons including improper use of the PMCE instances or improper calling of the (de-)serialization methods. No documentation could be found about specifics on those specific topics, though.

For Python libraries, one reason why documentation was in some cases sparse, only documented high-level features and largely omitted in-code documentation (such as comments) might be the “pythonic” approach to writing Python code. While

this way of programming may help to understand individual methods or even algorithms that use multiple methods, it does not by itself aid in documentation of high-level concepts or complex interaction. Another reason for sparse documentation in open source libraries might be the developers' focus on implementing more features or improving the code-base instead of aiming for more complete documentation. Contrary to commercial products, there usually are no monetary incentives for developers of open source software to write documentation.

### 8.2. Qualitative Overview: Deliverables

The following paragraphs evaluate what requirements were fulfilled by the design concepts discussed in chapter 6 and the implementation of the monolithic design presented in chapter 7.

**F1 Protocols:** Both design concepts address this requirement and share the same *IEncoder* interface that abstracts message (de-)serialization. Also, net-riot implemented most of the mandated suite of protocols (HTTP, WS, MQTT), lacking only an implementation of the Modbus TCP protocol.

**F2 Network Stacks:** The design concepts fulfil this requirement by envisioning a variant of the abstract factory pattern in the *Site* component used for dynamic instantiation of configured network stacks and FSMs. In net-riot, JSON and JSON schemas are used as the file format for configuration files, allowing formal specification of network stacks and FSMs.

**F3 State-Machines:** The monolithic design concept describes FSMs that feature an active states and rule-based transitions between states. The prototype implements the rules as scripts using Python's *exec* function.

**F4 Integration:** While neither of the design concepts defined explicit components or classes for integrating external software, they provide a basis that such integration components can be built upon (e.g. by deriving from *BasePipe*). However, the interfaces of such components need to be well-defined to ensure reusability. Consequently, net-riot did not implement integration capabilities.

**F5 Scripting:** The monolithic design concept fulfils this requirement by the definition of the *ProcessingPipe* class and the *IProcessor* interface that allow script-based transformation of messages. net-riot implements these classes and makes use of the *Script* class that provides execution of scripts at runtime by making use of Python's *exec* function.

**F6 Logging:** Similarly to "F4 Integration", logging capabilities were not explicitly defined by the design concepts and not implemented in net-riot.

**N1 Platform Compatibility:** Both design concepts are platform agnostic. Specifically, the distributed proxy design further improved on this requirement by allowing for distributing protocol implementations (i.e. *IEncoder* implementations) and transformation features (i.e. *IProcessor* implementations) across multiple machines. net-riot was written in Python, was tested on Ubuntu and should theoretically be executable on Windows and macOS machines, too.

**N2 Reusability:** By fulfilling "F2 Network Stacks" and "F3 State-Machines" and aiming for high flexibility in implementation, both design concepts promise high reusability. Due to its limitation on a single scenario and partial implementation, net-riot's reusability is limited.

**N3 Open Source:** The design concepts were conceived with being used in an open source project in mind. Their documentation (this thesis) provides high-level descriptions of their design principles and components. While the monolithic design is well-defined, the distributed design is not yet fully defined and requires further work (as described in section 6.2.4). As of writing this thesis, net-riot developed as an internal project at NVISO and therefore closed-source software. However, the goal is to make it available to the general public as open-source software once it reached a satisfying maturity level.

**N4 Extensibility:** Both design concepts feature high extensibility by hiding implementation details behind interface definitions. Protocol implementations can be implemented by implementing a single interface, *IEncoder*. Since net-riot defines the same interfaces and implements those itself, it too is highly extensible.

**N5 Deployment:** The design concepts do not address deployment explicitly. By choosing Python as its programming language though, net-riot can be packaged and containerized and thus employs repeatable and reliable methods for deployment.

### 8.3. Lessons Learned

During the various phases of the work performed during this thesis, a set of valuable lessons learned was found:

**Careful evaluation of third-party technologies:** shortcomings of software libraries such as restrictive design choices and a lack of documentation were recurring problems during implementation. Therefore, the design architecture and other documentation of such libraries should be studied before using them. Lack of sufficient documentation should be declared an exclusion criteria.

**Detect and react to an increase of problems during implementation:** Although the above-mentioned prior evaluation of third-party technologies may constitute an effective mitigation to the risk of said technologies failing, unforeseen problems can still occur during implementation. Thus, a reactive measure must be defined to mitigate this risk. When confronted with an ever-increasing number of problems, the technology choice should be re-evaluated as soon as possible.

**Conduct concise feasibility studies before committing to a project:** Especially in software projects that feature many and complex functionalities, concise feasibility studies should be conducted prior to committing to a project. These feasibility studies can be conducted by implementing Proof-of-Concept (PoC) prototypes and reviewing literature. They should be conducted in a well-defined time-span and evaluated afterwards. This can aid in identifying poorly understood requirements and infeasible technical requirements early.

**Define and adhere to an incremental development process:** The work on this thesis, the design concepts and implementations was performed in an iterative development process. However, it lacked clear evaluation of artefacts and was focused on implementation, neglecting necessary changes to the underlying design

concepts. Therefore, defining and adhering to an incremental development process is suggested to ensure regular evaluation of artefacts and design decisions.



## Chapter 9

### Summary

This chapter provides a summary of the design concepts presented in chapter 6 and the implementation of one of the concepts in chapter 7.

#### 9.1. Design Concepts

In chapter 6, the rough and vague design of the first prototype discussed in section 5.1.3 was further refined over the course of two iterations.

The first design concept presented in section 6.1 describes a monolithic proxy application that builds on the basic initial design which makes extensive use of the pipes and filters design pattern for message routing. At its core it sends and receives packets (*messages*) through gateways and (de-)serializes and transforms them through network stacks. It also features state-machines that allow implementation of complex logic of message (de-)serialization and transformation by allowing the binding of individual network stacks to states and triggering state transitions programmatically through the use of scripts. This results in a potentially deeply nested hierarchy of state-machines and network stacks.

The second design concept discussed in section 6.2 is an iteration of the monolithic design concept and describes a distributed proxy application that isolates (de-)serialization and transformation of messages from the internal proxy application logic. To decouple these low-level tasks from the high-level application logic of the proxy application, the concept introduces interfaces for remote units that provide access to specific features (i.e. (de-)serialization) and an interface for the central proxy application that allowed registering these remote units. Also, the nested hier-

archy of state-machines and network stacks is flattened and organized centrally in a repository in the central proxy application.

Both concepts feature distinct advantages and disadvantages and are compared in section 6.2.5.

### 9.2. Implementation

Chapter 7 shows an exemplaric implementation of the monolithic design concept discussed in section 6.1 under the working title “net-riot”. The monolithic design concept was chosen for implementation due to its proven core ideas and its comparatively high maturity. The second example scenario presented in section 5.1.1 was used for reference because it features a nested communication stack and a corresponding testbed has already been implemented in section 5.1.4. For this implementation, Python is used because of its flexibility, low barrier of entry and rich package ecosystem.

Since the reference scenario makes use of the HTTP, WS and MQTT protocols that used TCP as an underlying transport protocol, TCP gateways are implemented in net-riot as a MITM interface that external devices such as IoT devices and cloud server can connect to. For HTTP (de-)serialization, net-riot implements a custom encoder while for WS and MQTT existing libraries are used.

For representation of stacked communication protocols (such as MQTT being transported via WS), network stacks and state-machines were implemented: network stacks bundle a series of connected pipes that perform operations on messages, such as (de-)serializing and manipulating messages. State-machines allow selecting which network stacks to actively use by binding them to individual states. State-machines regularly evaluate their context and check whether states should be changed dependent on their registered transitions’ rules. These ScriptRules execute scripts which can examine and manipulate the states’ and state-machines’ context information.

A central task left open for implementation by the design concept is the configuration of the proxy application for specific scenarios and the resulting dynamic instantiation and parametrisation of state-machines and network stacks. In net-riot, JSON files and schemas were used for configuration specification and validation. Also, a recursive variance of the abstract factory design pattern was implemented for dynamic instantiation of objects defined in the configuration files. Figure 9.1



shows the output of the “cloc” utility program executed on net-riot’s source code. While net-riot implemented all components required by the scenario it was designed for, bugs in the implementation and challenging debugging ultimately lead to net-riot being unable to operate correctly in the testbed.

```
mo@ubuntu-vm:~/net-riot cloc.  
52 text files.  
52 unique files.  
3 files ignored.  
  
github.com/AlDanial/cloc v 1.82 T=0.04 s (1200.8 files/s, 115447.7 lines/s)  
-----  
Language          files      blank      comment      code  
-----  
Python             42         768         196         2744  
JSON                5           0           0         1045  
Markdown            1           2           0           42  
Bourne Shell        2           0           0           10  
-----  
SUM:                50         770         196         3841  
-----
```

**Figure 9.1.:** The lines of code in net-riot calculated by the “cloc” utility program.



## Chapter 10

### Conclusion

The goal of this thesis was to develop a concept for a modular proxy application that allows penetration testers to assess the security implementations of IoT applications.

Two design concepts were proposed; one featuring a monolithic and the other featuring a distributed software architecture. They share common components and interfaces for essential tasks such as protocol specific (de-)serialization of network packets. Due to these concepts being operating system, platform and framework agnostic, they can be implemented for a wide range of systems.

An exemplaric implementation of the monolithic design concept, net-riot, realized core components used for routing, (de-)serializing and transforming HTTP, WS and MQTT packets. Due to bugs in its implementation that affect the stacking of its MQTT and WS protocol implementations and the lack of time to fix these, net-riot is not operable in the scenario it was designed for yet. However, tests conducted with network stacks incorporating HTTP and WS communication were successful. While these bugs can most likely be resolved, future effort might instead be better invested in completing the distributed design concept and basing an implementation on it.

Especially the complex runtime behaviour and high amount of abstraction required to design protocol-agnostic interfaces proved to be challenges during the work on this thesis. Therefore, future work on this topic should aim to reduce the impact of these challenges by taking them into account during the development process:

- Flat data-structures and hierarchies can improve the traceability of data flow and thus support debugging.

- A contact person that is proficient in penetration testing and familiar with its challenges and requirements should be interviewed and asked for feedback on a regular basis.
- The creation and configuration of test environments should be streamlined so that testing can be performed regularly, ensuring that implemented features work correctly and according to specification.

### 10.1. Outlook

While the work on this thesis is completed, the project of designing and implementing a modular proxy application for testing IoT applications is not. There is a set of opportunities to continue this work:

- The distributed design concept promises attractive quality attributes attributes such as even better deployment and capabilities and extensibility. Also, its flattened hierarchy of FSMs and network stacks improve the debugging process. However, it was not fully defined and further work is required to fully define its components and interfaces.
- Consequently, an implementation based on the distributed design concept promises to be more feasible than net-riot's implementation based on the monolithic concept. The barrier of entry to such an implementation is lowered further due to the fact that the separation of proxy application and extensions allows free choice of multiple programming languages, platforms and frameworks for the systems.
- Of course, an evaluation of the usefulness of the modular proxy application is still an interesting endeavour. Whether it is based on the monolithic or the distributed design concept is not relevant since, from a black-box perspective, they perform the same tasks. Therefore future work could either begin at finishing net-riot or finalizing the distributed design concept.

# List of Abbreviations

|              |  |
|--------------|--|
| <b>A/C</b>   | Air Conditioner  |
| <b>ACSC</b>  | Australian Cyber Security Centre                                     |
| <b>API</b>   | application programming interface                                    |
| <b>ARP</b>   | Address Resolution Protocol  |
| <b>ASVS</b>  | Application Security Verification Standard                           |
| <b>AWS</b>   | Amazon Web Services  |
| <b>CERN</b>  | European Organization for Nuclear Research                           |
| <b>CR</b>    | Carriage Return  |
| <b>ENISA</b> | European Union Agency for Cybersecurity                              |
| <b>FSM</b>   | Finite-State Machine   |
| <b>GDPR</b>  | General Data Protection Regulation                                   |
| <b>GUI</b>   | Graphical User Interface   |
| <b>HMI</b>   | Human-Machine Interface  |
| <b>HTTP</b>  | Hypertext Transfer Protocol  |
| <b>HTTPS</b> | Hypertext Transfer Protocol Secure                                   |
| <b>ICS</b>   | Industrial Control System  |
| <b>IDE</b>   | Integrated Development Environment                                   |
| <b>IEC</b>   | International Electrotechnical Commission                            |
| <b>IIoT</b>  | Industrial Internet of Things  |
| <b>IoT</b>   | Internet of Things   |
| <b>IP</b>    | Internet Protocol  |
| <b>IPC</b>   | Inter-Process Communication  |
| <b>ISO</b>   | International Organization for Standardization                       |
| <b>ISP</b>   | Internet Service Provider  |
| <b>ISVS</b>  | IoT Security Verification Standard                                   |
| <b>JSON</b>  | JavaScript Object Notation   |
| <b>LAN</b>   | Local Area Network   |
| <b>LF</b>    | Line Feed  |
| <b>MAC</b>   | Media Access Control   |
| <b>MITM</b>  | Man-in-The-Middle  |
| <b>MQTT</b>  | Message Queuing Telemetry Transport                                  |
| <b>MSVS</b>  | Mobile Application Security Verification Standard                    |
| <b>MTU</b>   | Maximum Transmission Unit  |
| <b>OASIS</b> | Organization for the Advancement of Structured Information Standards |

|                |  |
|----------------|--|
| <b>OPC U/A</b> | OPC Unified Architecture               |
| <b>OSI</b>     | Open Systems Interconnection           |
| <b>OSINT</b>   | Open Source Intelligence               |
| <b>OSS</b>     | Open Source Software                   |
| <b>OWASP</b>   | Open Web Application Security Project  |
| <b>PLC</b>     | Programmable Logic Controller          |
| <b>PMCE</b>    | Per-Message Compression Extension      |
| <b>PoC</b>     | Proof-of-Concept                       |
| <b>PTES</b>    | Penetration Testing Execution Standard |
| <b>PyPI</b>    | Python Package Index                   |
| <b>QoS</b>     | Quality of Service                     |
| <b>REST</b>    | Representational State Transfer        |
| <b>RPC</b>     | Remote Procedure Call                  |
| <b>SSL</b>     | Secure Socket Layer                    |
| <b>TCP</b>     | Transmission Control Protocol          |
| <b>TLS</b>     | Transport Layer Security               |
| <b>UDP</b>     | User Datagram Protocol                 |
| <b>USB</b>     | Universal Serial Bus                   |
| <b>WLAN</b>    | Wireless LAN                           |
| <b>WS</b>      | WebSocket                              |
| <b>XML</b>     | Extensible Markup Language             |

## List of Figures

|       |  |    |
|-------|--|----|
| 3.1.  | An abstract sequence diagram of the pipes and filters design pattern. The relaying mechanism used by pipes in this diagram is realized through method calls. . . . .   | 6  |
| 3.2.  | “Seven layer reference model and peer protocols” [18] as proposed in ISO/IEC 7498. . . . .   | 8  |
| 3.3.  | A truncated HTTP request (indicated by “>”) to “httpbin.org” using the utility “curl” and the truncated response received from the remote server (indicated by “<”). . . . .   | 11 |
| 3.4.  | Screenshot of Wireshark being executed and dissecting a HTTP <i>GET</i> request to the site “httpbin.org”. The display-filter “tcp.port == 80” shows only packets sent to or from port 80 (e.g. HTTP communication). . . . .                           | 16 |
| 3.5.  | Screenshot of Burp Suite being used to send forged HTTP requests to the site “httpbin.org”. . . . .  | 17 |
| 5.1.  | Installing a MITM proxy to intercept network communication for penetration testing. . . . .  | 22 |
| 5.2.  | The simplistic architecture of the first scenario. . . . .   | 22 |
| 5.3.  | The central MQTT broker relays messages to subscribed clients. In this case, the A/C unit receives messages published to a topic it subscribed to. . . . .   | 23 |
| 5.4.  | State machine of AWS IoT communication . . . . .   | 24 |
| 5.5.  | Illustration of a typical drinking water treatment process. (by the CK-12 Foundation) . . . . .  | 25 |
| 5.6.  | MQTT clients for each station of the water treatment plant publish messages about their tanks’ levels to individual topics and subscribe to their succeeding station’s tank levels in order to simulate flushing water into their input tanks. . . . . | 26 |
| 5.7.  | High-level use-cases of a proxy in a generic IoT/ICS environment. . . . .  | 27 |
| 5.8.  | The variation of the “pipes and filters” design pattern used in the prototype. . . . .   | 30 |
| 5.9.  | AWS IoT Scenario - State 1: HTTP Server . . . . .  | 31 |
| 5.10. | AWS IoT Scenario - State 2: MQTT via WS . . . . .  | 33 |

|  |    |
|--|----|
| 5.11. A network diagram of the testbed that was used for testing the prototype. . . . .  | 34 |
| 5.12. The “ProcessingUnit” data-structures represent individual stations of the simplified water treatment plant. . . . .  | 34 |
| 5.13. Chaining of the water treatment units, originating from a water source and eventually leading to a storage at the end of the processing pipeline. . . . .  | 35 |
| 5.14. Screenshot of the application “MQTT Explorer” that was used to inspect and visualize the state of the water treatment plant. The left graph shows how the <i>source</i> ’s input tank steadily emptied until it was filled by the <i>storage</i> ’s output tank. The right graph shows how the <i>flocculant</i> unit’s input tank slowly filled up. . . . . | 35 |
| 5.15. The classes and interfaces used to implement pipelines in the TypeScript prototype. . . . .  | 36 |
| 5.16. The structure of messages in the prototype: messages hold information about their communication direction, protocol-specific meta data and type-dependent payloads. . . . .  | 37 |
| 5.17. A variation of the abstract factory design pattern was used to decouple the creation of objects from their usage. . . . .  | 38 |
| 6.1. High-level component diagram of the proxy application concept . .   | 47 |
| 6.2. A simple variation of the abstract factory pattern. Contrary to the design of the first prototype in section 5.1.3, this variation does not specify the abstract type of the products that are built as return types but as part of the meta data used for object creation. . . . .   | 48 |
| 6.3. The “Message IO”-package. The white boxes indicate ports for communication with outside components. . . . .   | 48 |
| 6.4. An abstract representation of the runtime hierarchy of nested FSMs and network stacks. The active states of the FSMs determine which network stacks are used, resulting in a chain of FSMs and network stacks, referred to as a “pipeline”. . . . .   | 49 |
| 6.5. Message processing through an architecture of nested FSMs and network stacks . . . . .  | 50 |
| 6.6. <i>StateMachines</i> used rules to determine whether state changes should take place. <i>IChainRules</i> , <i>AndRules</i> and <i>OrRules</i> allowed to combine multiple rules and provided a way to specify logic using <i>OR</i> and <i>AND</i> operators. . . . .   | 51 |
| 6.7. By introducing <i>Gateways</i> , <i>GatewayPipes</i> and <i>RelayPipes</i> , multiplexing pipes could be implemented. . . . .   | 52 |
| 6.8. The interfaces, classes and enums used to represent pipes, their specializations and associated classes. This iteration separates <i>BasePipes</i> from <i>IEncoders</i> and <i>IProcessors</i> so <i>BasePipes</i> only implement routing of messages. . . . .   | 53 |
| 6.9. The data-structures used to represent messages, their meta-data and payloads. . . . .   | 55 |



|       |   |      |
|-------|---|------|
| 6.10. | A component diagram showing the external systems communicating with the system and their connection to internal components. . .   | 56   |
| 6.11. | In this concept, $n$ separate <i>Units</i> can be registered at a central <i>Registry</i> . <i>Units</i> may provide an arbitrary amount of <i>Features</i> . . . . .   | 57   |
| 6.12. | Visualized in this diagram, the distributed implementation and central registration of features is the core idea of the Registry and Unit components. . . . .   | 58   |
| 7.1.  | The implementation of the variant of the abstract factory pattern discussed in section 6.1. The Site component and IFactory instances recursively call each other to create and composite object instances.                         | 67   |
| 9.1.  | The lines of code in net-riot calculated by the “cloc” utility program.   | 85   |
| A.1.  | The interfaces, classes and enums used to represent pipes, their specializations and associated classes. This iteration separates BasePipes from IEncoders and IProcessors so BasePipes only implement routing of messages. . . . . | xxvi |



## List of Tables

|  |    |
|--|----|
| 5.1. Comparison of existing software . . . . .           | 46 |
| 8.1. Initially planned schedule for the thesis . . . . . | 72 |
| 8.2. Actual schedule of the project . . . . .            | 74 |



## Listings

|  |    |
|--|----|
| 7.1. The ScriptProcessor implementation of net-riot using Python’s built-in exec function. . . . .           | 66 |
| 7.2. The script net-riot used to detect upgrades of HTTP connections to WS. . . . .                          | 66 |
| 7.3. A simple factory that requested the build of a “Script” instance to instantiate a “ScriptRule”. . . . . | 68 |
| 7.4. Dynamic registering of all factory implementations. . . . .   | 68 |



# Bibliography

- [1] P. S. C. Alencar, D. D. Cowan, and C. J. P. Lucena. “A formal approach to architectural design patterns”. In: *FME96: Industrial Benefit and Advances in Formal Methods Lecture Notes in Computer Science* (1996), pp. 576–594. DOI: 10.1007/3-540-60973-3\_108.
- [2] S. Ansari, S.G. Rajeev, and H.S. Chandrashekar. “Packet sniffing: a brief introduction”. In: *IEEE Potentials* 21.5 (2003), pp. 17–19. DOI: 10.1109/MP.2002.1166620.
- [3] Noah Apthorpe, Dillon Reisman, and Nick Feamster. *A Smart Home is No Castle: Privacy Vulnerabilities of Encrypted IoT Traffic*. 2017. arXiv: 1705.06805 [cs.CR].
- [4] Michael Bartsch et al. *Spionage, Sabotage und Datendiebstahl - Wirtschaftsschutz in der Industrie 2018*. Oct. 2018. URL: <https://www.bitkom.org/sites/default/files/file/import/181008-Bitkom-Studie-Wirtschaftsschutz-2018-NEU.pdf>.
- [5] Cédric Bassem. *Expert Interview*. Apr. 2020.
- [6] Jonah Bellemans. *Expert Interview*. July 2020.
- [7] Jonah Bellemans. “The state of the market: A comparative study of IoT device security implementations”. MA thesis. KU Leuven, 2020.
- [8] *bettercap*. URL: <https://www.bettercap.org/> (visited on 05/22/2021).
- [9] Jimmy Braden. *Penetration Testing - Is it right for you?* 2002. URL: <https://www.sans.org/reading-room/whitepapers/testing/penetration-testing-you-265>.
- [10] Dan Demeter, Marco Preuss, and Yaroslav Shmelev. *IoT: a malware story*. Oct. 2019. URL: <https://securelist.com/iot-a-malware-story/94451/> (visited on 02/17/2020).
- [11] Patrick Eisenschmidt. *Expert Interview*. Apr. 2020.

- [12] *ettercap*. URL: <https://github.com/Ettercap/ettercap/> (visited on 05/22/2021).
- [13] Erich Gamma et al. *Design Patterns - Elements of Reusable Object-oriented Software*. Pearson Education, 1994.
- [14] Rahul Gupta et al. *MQTT Version 5.0*. Mar. 2019. URL: <https://docs.oasis-open.org/mqtt/mqtt/v5.0/mqtt-v5.0.html> (visited on 05/28/2021).
- [15] Simon Hansman and Ray Hunt. “A taxonomy of network and computer attacks”. In: *Computers & Security* 24.1 (2005), pp. 31–43. DOI: <https://doi.org/10.1016/j.cose.2004.06.011>. URL: <https://www.sciencedirect.com/science/article/pii/S0167404804001804>.
- [16] *IOXY - MQTT intercepting proxy*. URL: <https://github.com/NVISOsecurity/IOXY> (visited on 05/22/2021).
- [17] *Information technology — Security techniques — Information security management systems — Overview and vocabulary*. Standard. Geneva, CH: International Organization for Standardization, International Electrotechnical Commission, Feb. 2018.
- [18] *Information processing systems — Open Systems Interconnection — Basic Reference Model: The Basic Model*. Standard. Geneva, CH: International Organization for Standardization, International Electrotechnical Commission, Nov. 1994.
- [19] Jens Jäger et al. “Advanced Complexity Management Strategic Recommendations of Handling the “Industrie 4.0” Complexity for Small and Medium Enterprises”. In: *Procedia CIRP*. Factories of the Future in the digital environment - Proceedings of the 49th CIRP Conference on Manufacturing Systems 57 (Jan. 2016), pp. 116–121. DOI: 10.1016/j.procir.2016.11.021.
- [20] Christian Lesjak et al. “Security in industrial IoT – quo vadis?” In: *e & i Elektrotechnik und Informationstechnik* 133.7 (Nov. 2016), pp. 324–329. DOI: 10.1007/s00502-016-0428-4. URL: <https://doi.org/10.1007/s00502-016-0428-4>.
- [21] B. Lundgren and Niklas Möller. “Defining Information Security”. In: *Science and Engineering Ethics* 25 (2019), pp. 419–441.
- [22] Alexey Melnikov and Ian Fette. *The WebSocket Protocol*. RFC 6455. Dec. 2011. DOI: 10.17487/RFC6455. URL: <https://rfc-editor.org/rfc/rfc6455.txt>.

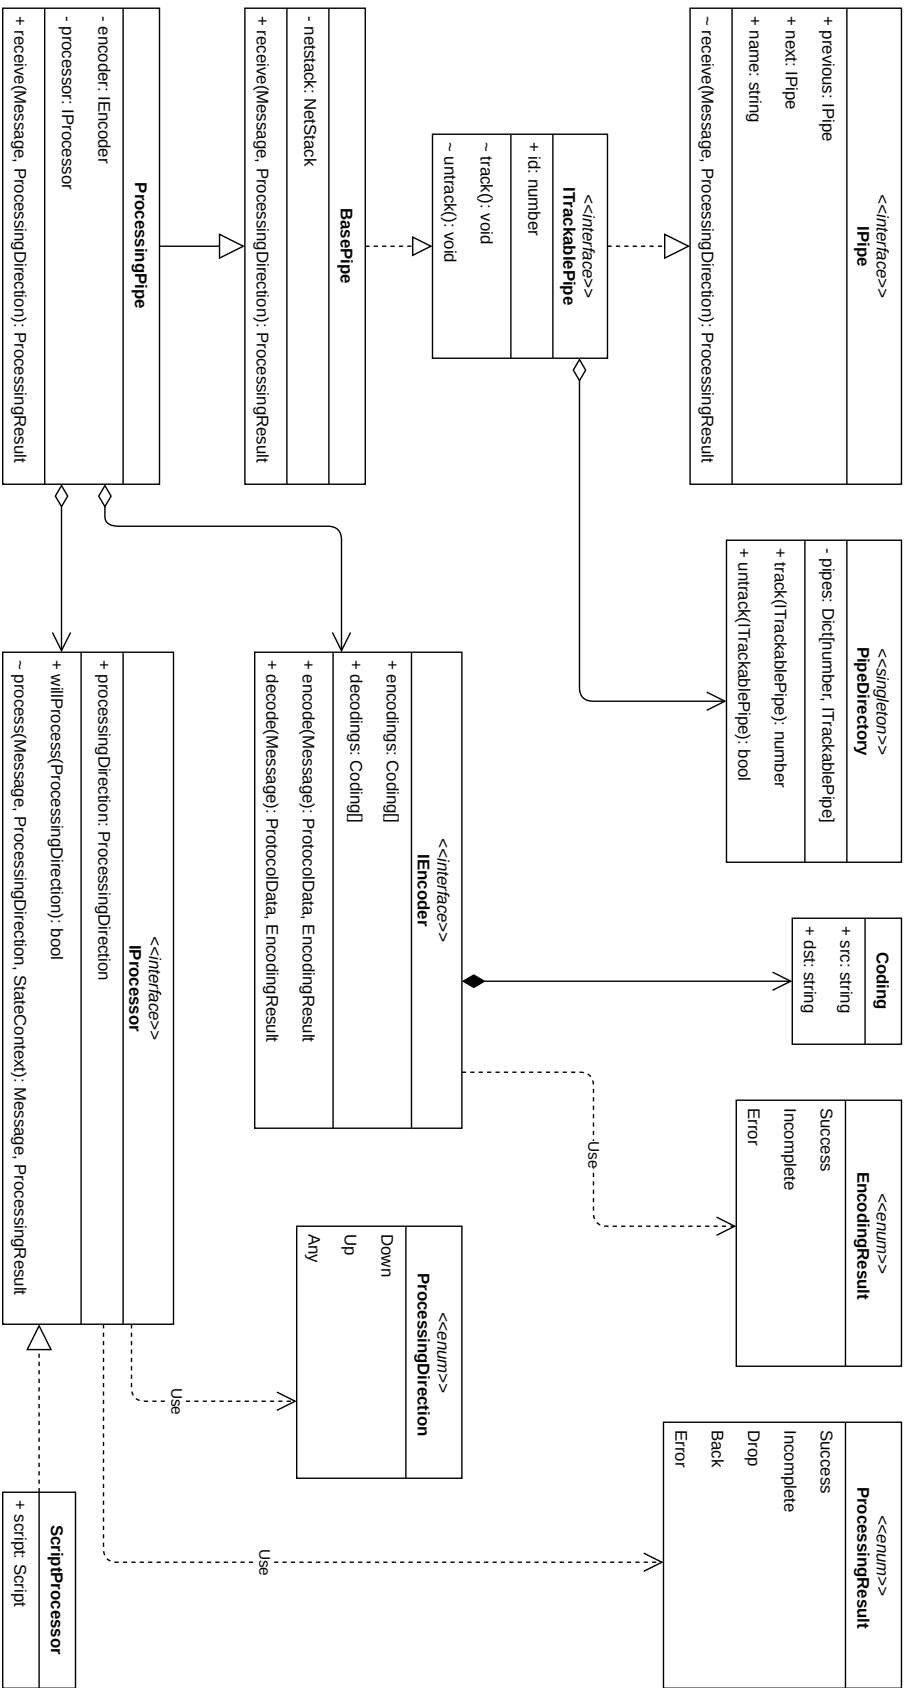


- 
- [23] *MITMf - Framework for Man-In-The-Middle attacks*. URL: <https://github.com/byt3bl33d3r/MITMf> (visited on 05/22/2021).
- [24] *mitmproxy*. URL: <https://mitmproxy.org/> (visited on 05/22/2021).
- [25] Pierre-Alain Mouy. *Expert Interview*. Apr. 2020.
- [26] *mProxy*. URL: <https://github.com/mainflux/mproxy> (visited on 05/22/2021).
- [27] Oliver Nettinger. *Expert Interview*. May 2020.
- [28] Henrik Nielsen, Roy T. Fielding, and Tim Berners-Lee. *Hypertext Transfer Protocol – HTTP/1.0*. RFC 1945. May 1996. DOI: 10.17487/RFC1945. URL: <https://rfc-editor.org/rfc/rfc1945.txt> (visited on 05/22/2021).
- [29] Jacob G. Oakley. *Professional Red Teaming - Conducting Successful Cybersecurity Engagements*. New York: Apress, 2019.
- [30] *OWASP Application Security Verification Standard*. URL: <https://github.com/OWASP/ASVS> (visited on 05/28/2021).
- [31] *OWASP IoT Security Verification Standard*. URL: <https://github.com/OWASP/IoT-Security-Verification-Standard-ISVS> (visited on 05/28/2021).
- [32] *OWASP Mobile Application Security Verification Standard*. URL: <https://github.com/OWASP/owasp-masvs> (visited on 05/28/2021).
- [33] *OWASP Penetration Testing Execution Standard*. URL: [https://owasp.org/www-project-web-security-testing-guide/v41/3-The\\_OWASP\\_Testing\\_Framework/1-Penetration\\_Testing\\_Methodologies](https://owasp.org/www-project-web-security-testing-guide/v41/3-The_OWASP_Testing_Framework/1-Penetration_Testing_Methodologies) (visited on 05/29/2021).
- [34] Charith Perera, Chi Harold Liu, and Srimal Jayawardena. “The Emerging Internet of Things Marketplace From an Industrial Perspective: A Survey”. In: *IEEE Transactions on Emerging Topics in Computing* 3.4 (2015), pp. 585–598. DOI: 10.1109/TETC.2015.2390034.
- [35] *Person-in-the-middle*. May 2020. URL: <https://www.cyber.gov.au/acsc/view-all-content/glossary/person-middle>.
- [36] *PortSwigger - Burp Suite*. URL: <https://portswigger.net/burp> (visited on 05/22/2021).
- [37] Eric Rescorla. *HTTP Over TLS*. RFC 2818. May 2000. DOI: 10.17487/RFC2818. URL: <https://rfc-editor.org/rfc/rfc2818.txt> (visited on 05/22/2021).

- [38] Théo Rigas. *Expert Interview*. Apr. 2020.
- [39] *Scapy - Packet crafting for Python2 and Python3*. URL: <https://owasp.org/www-project-zap/> (visited on 05/22/2021).
- [40] Hashim Shaikh. *Pivoting to exploit a system in another network*. May 2017. URL: <https://resources.infosecinstitute.com/topic/pivoting-exploit-system-another-network/>.
- [41] *socat*. URL: <https://repo.or.cz/socat.git> (visited on 05/22/2021).
- [42] Michael Sutton, Adam Greene, and Pedram Amini. *Fuzzing - Brute Force Vulnerability Discovery*. Amsterdam: Pearson Education, 2007.
- [43] *The Original HTTP as defined in 1991*. URL: <https://www.w3.org/Protocols/HTTP/AsImplemented.html> (visited on 05/22/2021).
- [44] *Wireshark*. URL: <https://www.wireshark.org/> (visited on 05/22/2021).
- [45] *Wireshark - Git Repository (commit c73ab16b)*. URL: <https://gitlab.com/wireshark/wireshark/> (visited on 05/22/2021).
- [46] Tilman Wittenhorst. *Ferngesteuert ins Smart Home: Saugroboter verrät Grundriss der Wohnung*. June 2019. URL: <https://www.heise.de/security/meldung/Ferngesteuert-ins-Smart-Home-Saugroboter-verraet-Grundriss-der-Wohnung-4436657.html> (visited on 02/17/2020).
- [47] Takeshi Yoshino. *Compression Extensions for WebSocket*. RFC 7692. Dec. 2015. DOI: 10.17487/RFC7692. URL: <https://rfc-editor.org/rfc/rfc7692.txt>.

## **Appendix A**

### **Diagrams**



**Figure A.1:** The interfaces, classes and enums used to represent pipes, their specializations and associated classes. This iteration separates BasePipes from IEncoders and IProcessors so BasePipes only implement routing of messages.

## **Appendix B**

### **Interview Guideline**

#### **1. Experiences with IoT**

- a) Which technologies (software/protocols/platforms) were used in those applications?
- b) What context (home/industrial) were they used in?
- c) Were there any special constraints (e.g. real-time systems) when working with them?
- d) How were the tests set up? (e.g. dedicated lab vs. on-site testing, single/multiple devices)
- e) Were those applications typical representatives for their field of use?

#### **2. Processes in Everyday Life**

- a) What are the goals and scopes of your penetration tests?
- b) Which tasks do these goals typically involve?
- c) Which tools do you use in the process and how regularly do you create specific tools yourself?
- d) What problems do you typically face during tests?
- e) Could specialized tools further improve your everyday work? If so, what would those tools do?

#### **3. The Future of IoT**

- a) What do you think will IoT applications be like in the future? (e.g. regarding device specifications, cloud services, mobile integrations)
- b) What are the current challenges when working with IoT applications?
- c) Lastly, what do you think will be future challenges when working with IoT?